

# Generative Models I

Jisang Han

onground@korea.ac.kr

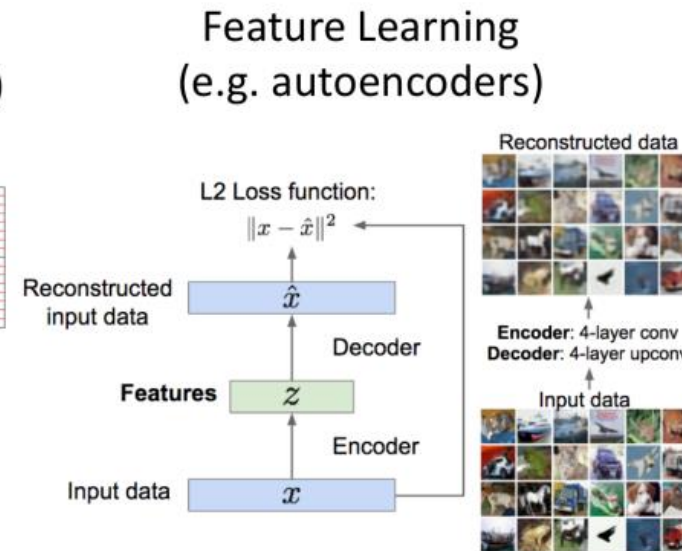
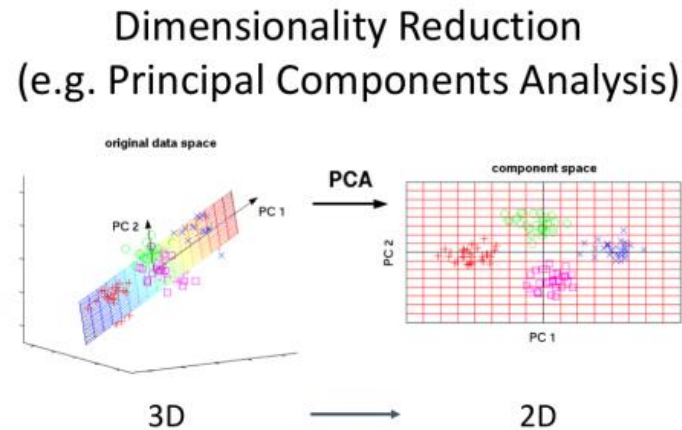
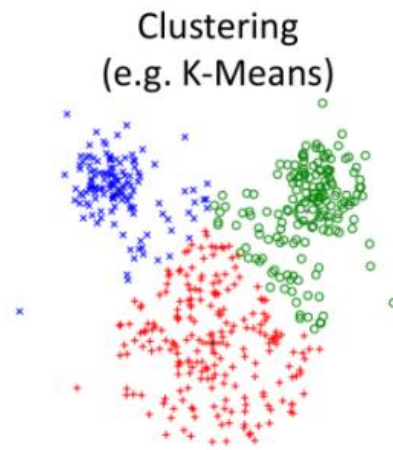
Artificial Intelligence in KU (AIKU)

Department of Computer Science and Engineering, Korea University

**AIKU**

# Supervised vs Unsupervised

	Supervised Learning	Unsupervised Learning
<b>Data</b>	$(x, y)$ : $x$ is data, $y$ is label	$x$
<b>Goal</b>	$x \rightarrow y$ 로 가는 function을 학습	Learn some underlying hidden structure of the data
<b>Exam ples</b>	Classification, regression, object detection, semantic segmentation, image captioning, etc.	Clustering, dimensionality reduction, feature learning, density estimation, etc.



# Discriminative vs Generative Models

---

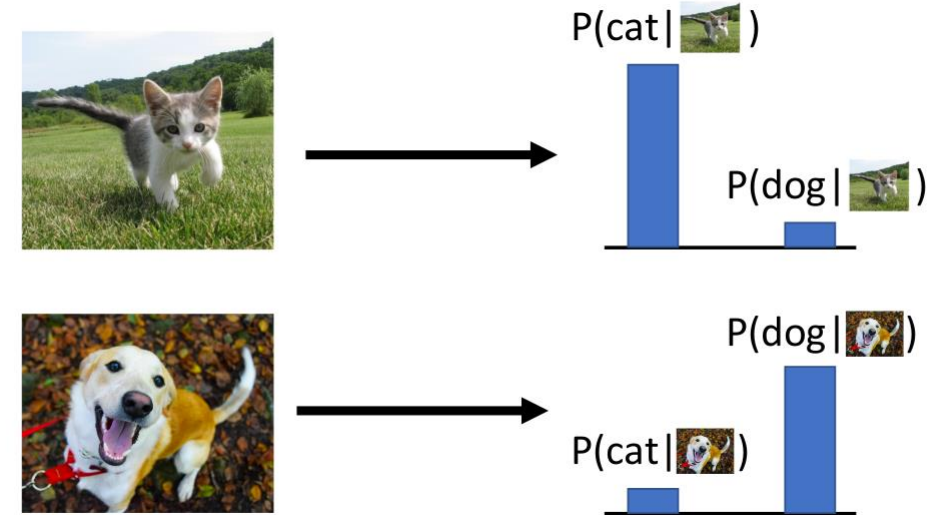
## Density Function

$$\int_X p(x) dx = 1$$

- Density Functions are normalized.
- Sum of probabilities are 1. If one is bigger then the other is smaller(competes each other).

# Discriminative vs Generative Models

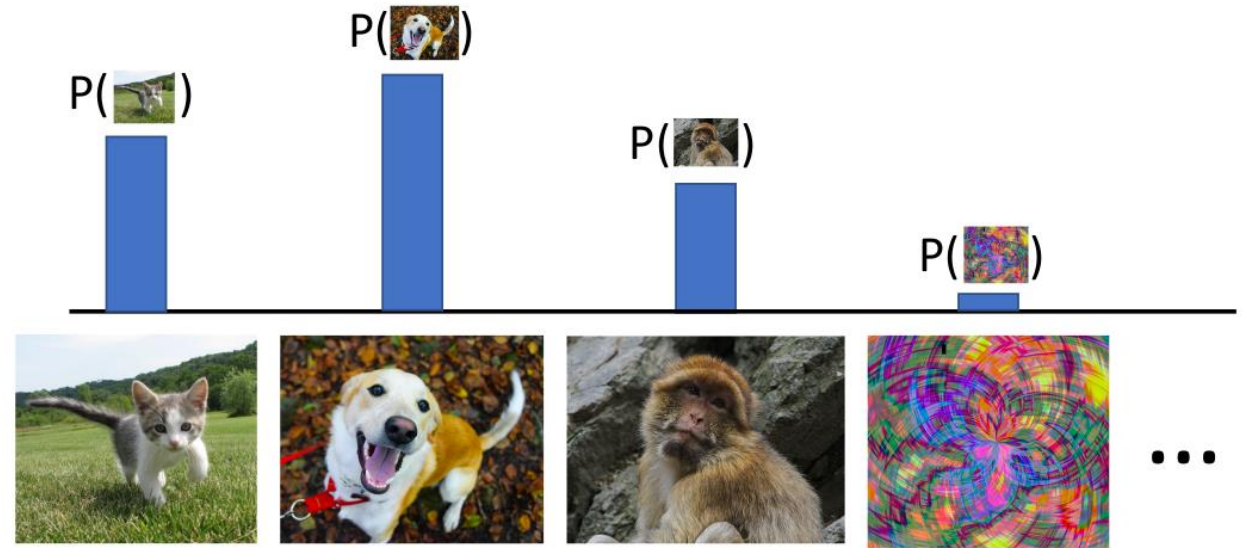
## Discriminative Model



- Learn a probability distribution  $p(y|x)$  that predicts probability of the label  $y$  conditioned on the input image  $x$
- Input is  $x$  and Output is the probability. ( $x$ 에 대한 label이 나올 확률)
- The probability is calculated even if an input different from the specified labels is given.

# Discriminative vs Generative Models

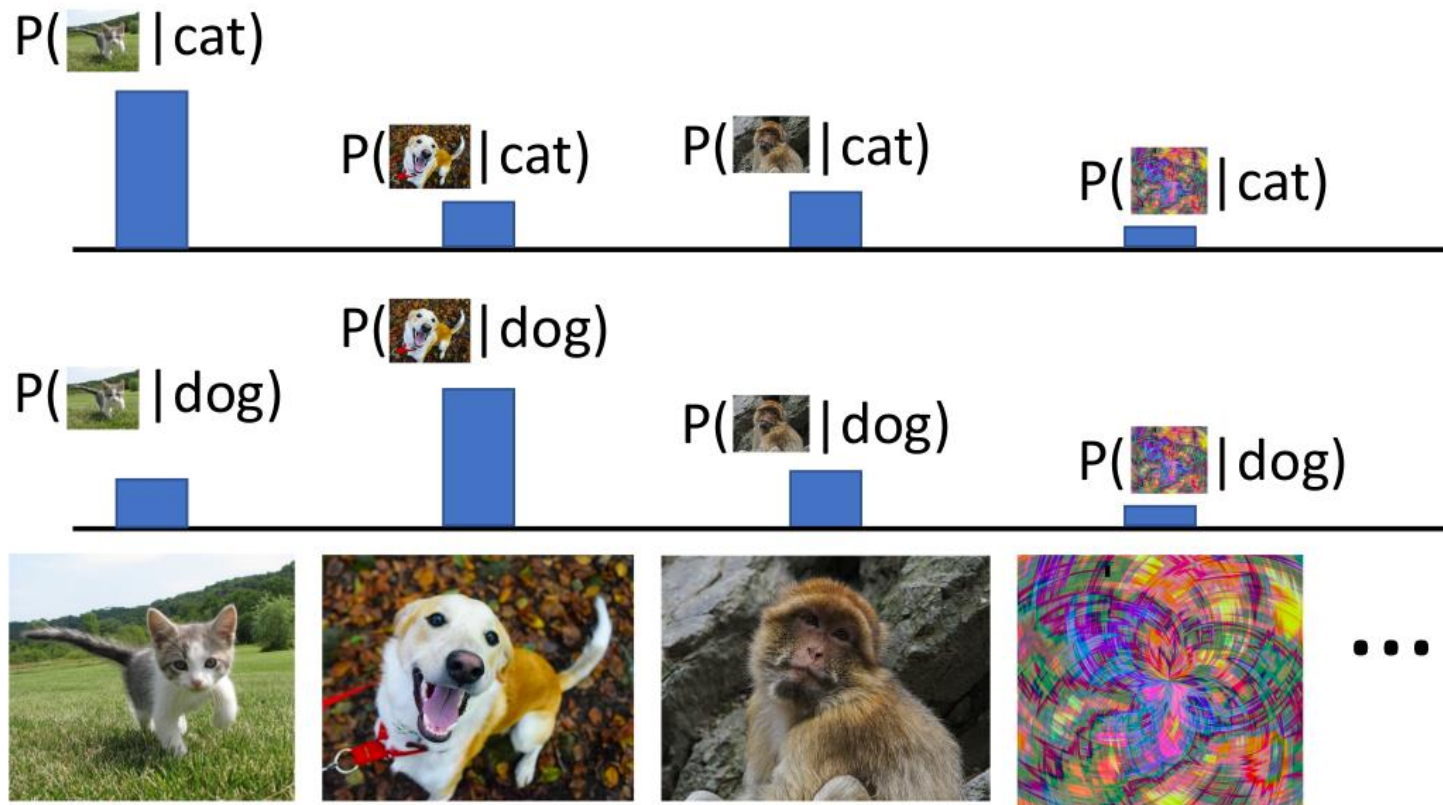
## Generative Model



- **Learn a probability distribution  $p(x)$**
- All possible images compete to the probability mass.
- A deep understanding of the image is needed. (Which is more plausible, sitting a dog or standing up?)
- The model can reject irrational inputs by giving them a very small value.

# Discriminative vs Generative Models

## Conditional Generative Model



Recall Bayes' Rule:

$$P(x | y) = \frac{P(y | x) P(x)}{P(y)}$$

Bayes' Rule equation with labels:

- $P(x | y)$ : Conditional Generative Model
- $P(y | x)$ : Discriminative Model
- $P(x)$ : (Unconditional) Generative Model
- $P(y)$ : Prior over labels

- Learn  $p(x | y)$

# Discriminative vs Generative Models

---

## **Discriminative Model:**

Learn a probability distribution  $p(y|x)$



Assign labels to data  
Feature learning (supervised)

## **Generative Model:**

Learn a probability distribution  $p(x)$



Detect outliers  
Feature learning (unsupervised)  
Sample to **generate** new data

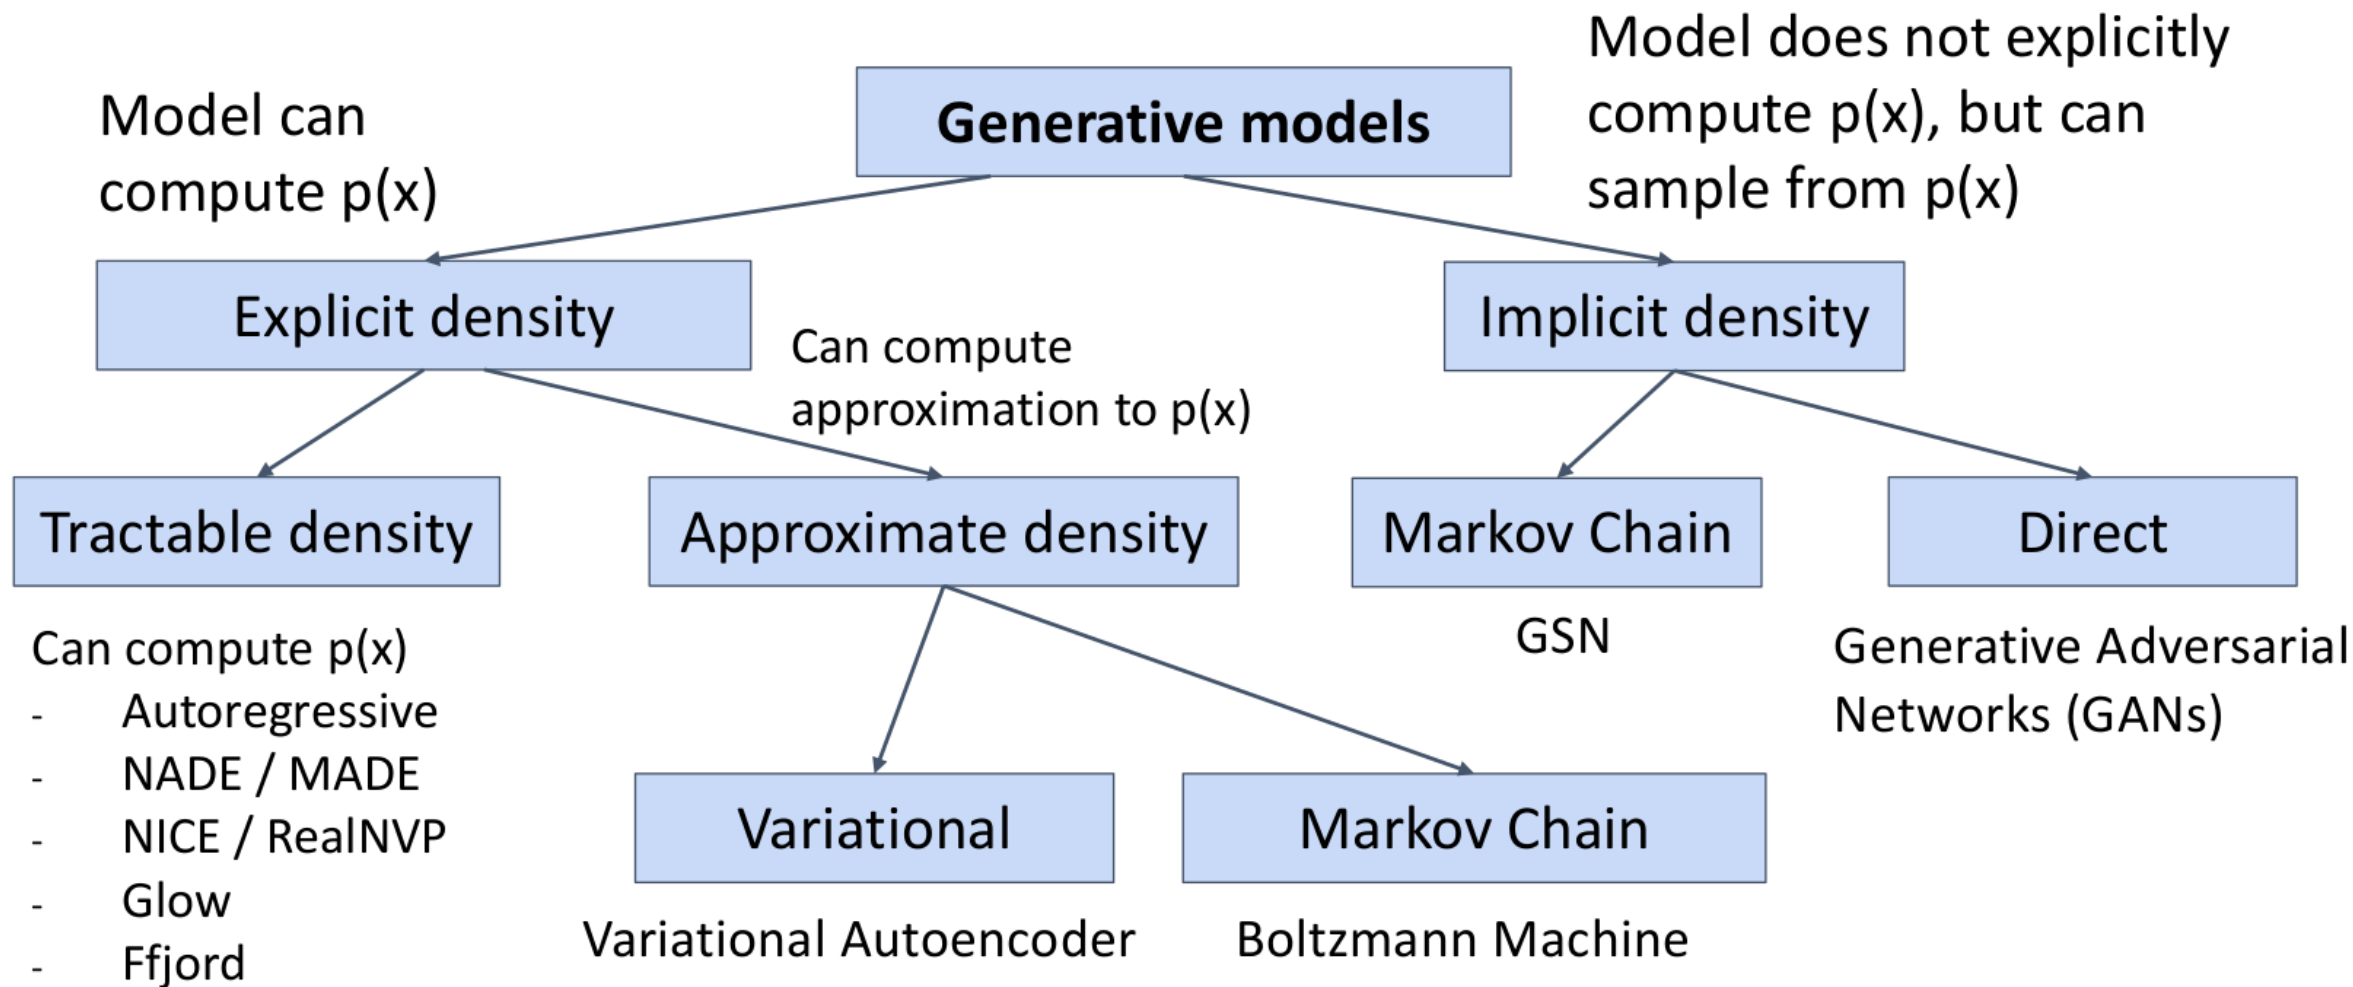
## **Conditional Generative Model:** Learn $p(x|y)$



Assign labels, while rejecting outliers!  
Generate new data conditioned on input labels



# Taxonomy of Generative Models





# Autoregressive Models

---

## Explicit Density Estimation

**Goal :** Write down an explicit function for  $p(x) = f(x, W)$  ( $x$ : data,  $W$ : learnable weight matrix)

If dataset is  $x^{(1)}, x^{(2)}, \dots, x^{(N)}$ ,

$$\begin{aligned} W^* &= \arg \max_W \prod_i p(x^{(i)}) \\ &= \arg \max_W \sum_i \log p(x^{(i)}) \\ &= \arg \max_W \sum_i \log f(x^{(i)}, W) \end{aligned}$$

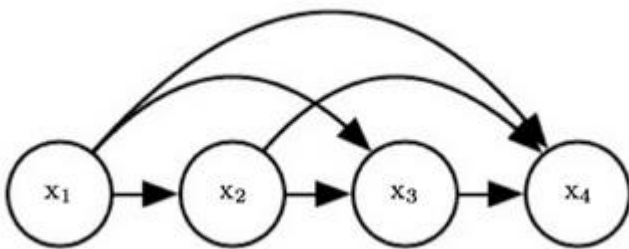
Maximize the probability of training data

Loss Function. (Gradient Descent)

# Autoregressive Models

## Explicit Density: Autoregressive Models

- 자기 자신을 입력으로 하여 자기 자신을 예측하는 모형



- $x$ 가 여러 subparts로 이루어져있다고 가정.  $x$ 가 이미지라고 하면 subparts는 각 픽셀이다  $x = (x_1, x_2, x_3, \dots, x_T)$
- Chain rule을 사용하여 probability를 계산한다.

$$\begin{aligned} p(x) &= p(x_1, x_2, \dots, x_T) \\ &= p(x_1) p(x_2|x_1) p(x_3|x_1, x_2) \\ &= \prod_{t=1}^T p(x_t|x_1, \dots, x_{t-1}) \end{aligned}$$

# Autoregressive Models

---

## Explicit Density: Autoregressive Models

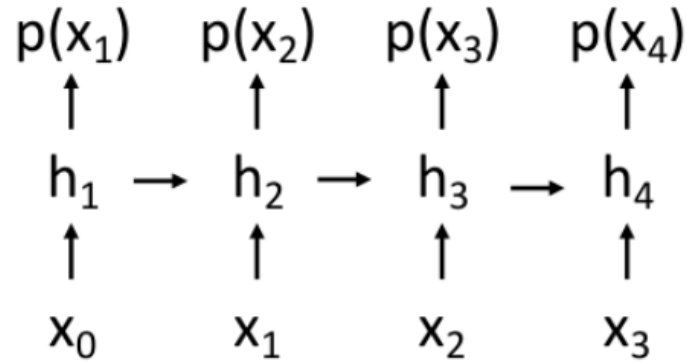
$$\begin{aligned} p(x) &= p(x_1, x_2, \dots, x_T) \\ &= p(x_1) p(x_2|x_1) p(x_3|x_1, x_2) \\ &= \prod_{t=1}^T p(x_t|x_1, \dots, x_{t-1}) \end{aligned}$$

- 즉, sequence의 확률  $p(x)$ 는 이전 sequence가 주어졌을 때 다음 sequence가 나올 확률을 전부 곱한 것이다.
- <https://wikidocs.net/22034>

# Autoregressive Models

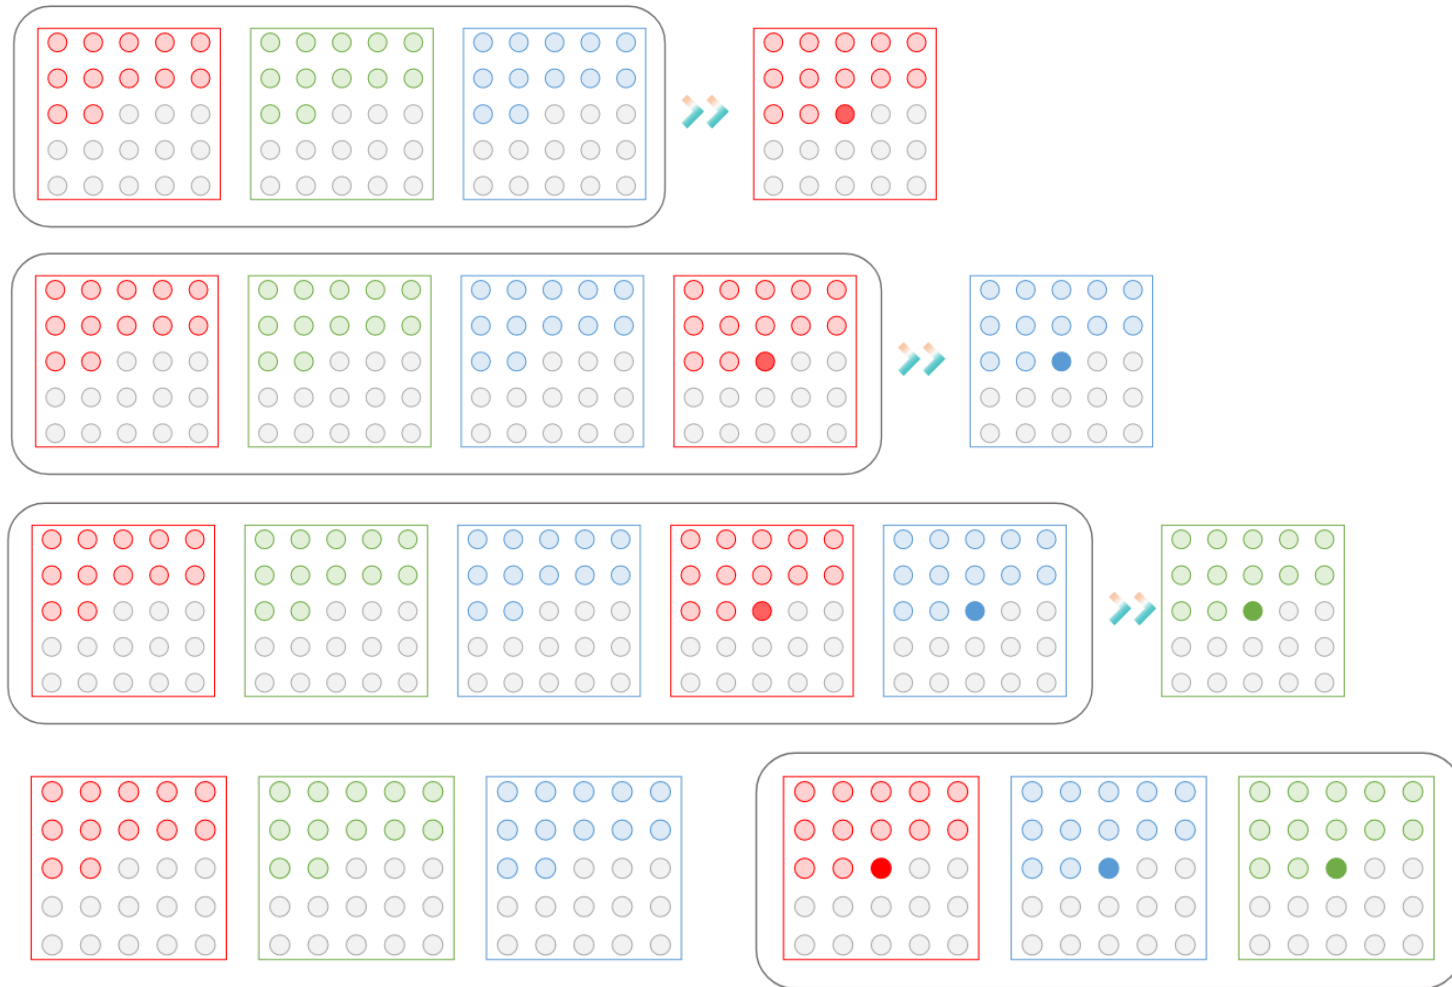
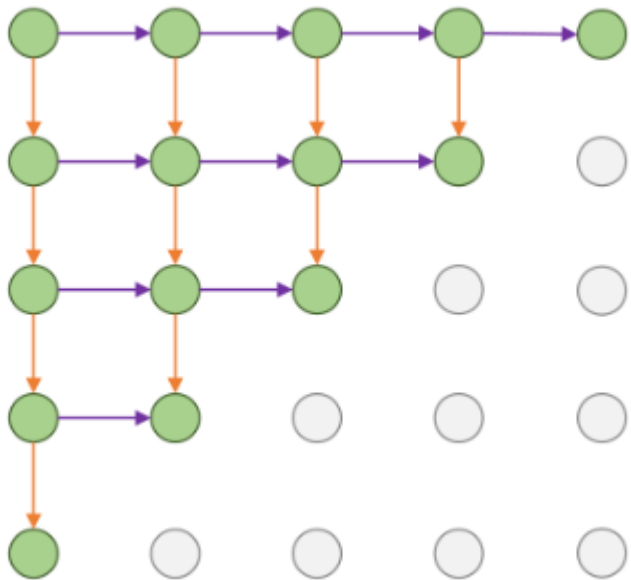
---

## Explicit Density: Autoregressive Models



# Autoregressive Models

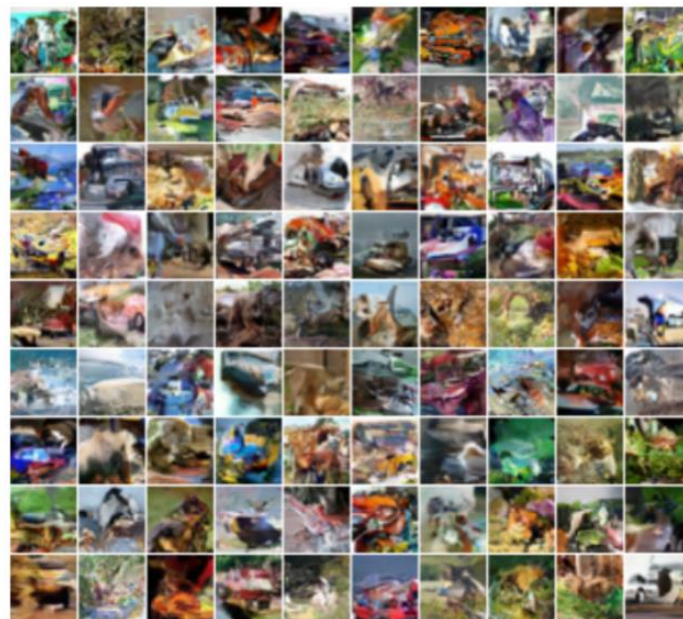
## PixelRNN



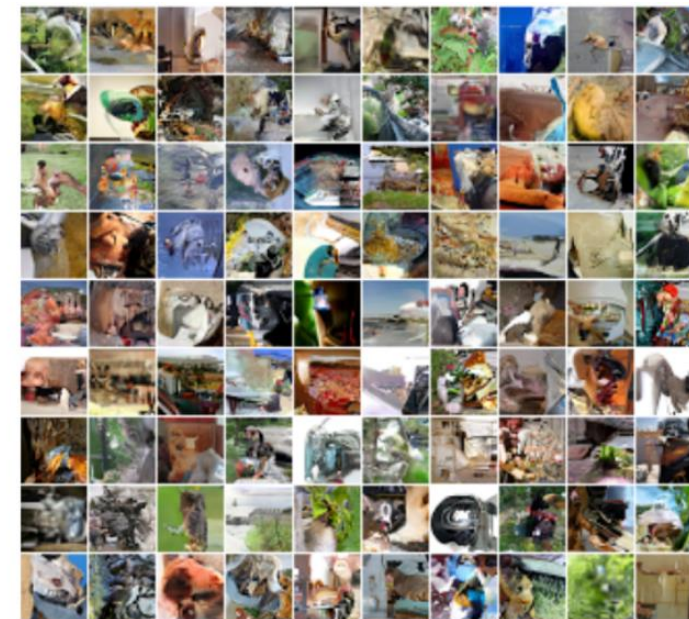
[http://dmqm.korea.ac.kr/uploads/seminar/20190705\\_Autoregressive.pdf](http://dmqm.korea.ac.kr/uploads/seminar/20190705_Autoregressive.pdf)

# Autoregressive Models

## PixelRNN



32x32 CIFAR-10



32x32 ImageNet

- 겉으로 보기에는 합리적으로 보이지만 사실 자세히 보면 거지같다..
- 그래도 edges, colors를 꽤나 그럴싸하게 생성하는 것에 의의를 가진다.
- 이는 unconditional generation이므로 test time에 내가 무엇을 생성하고 있는지 정할 수 없다

# Autoregressive Models

---

## Pros and Cons

### Pros

- Likelihood  $p(x)$ 를 명시적으로 계산할 수 있다.
- 위의 장점 덕분에 좋은 evaluation metric을 얻는다. (training data와 유사한, 즉 얼마나 그럴듯한지에 대한  $p(x)$  값을 얻을 수 있기 때문이다.)
- 나름 괜찮은 결과를 얻을 수 있다. (진짜 사진같지는 않지만..)

### Cons

- Sequential Generation 이므로 느리다.



# Variational Autoencoders (VAE)

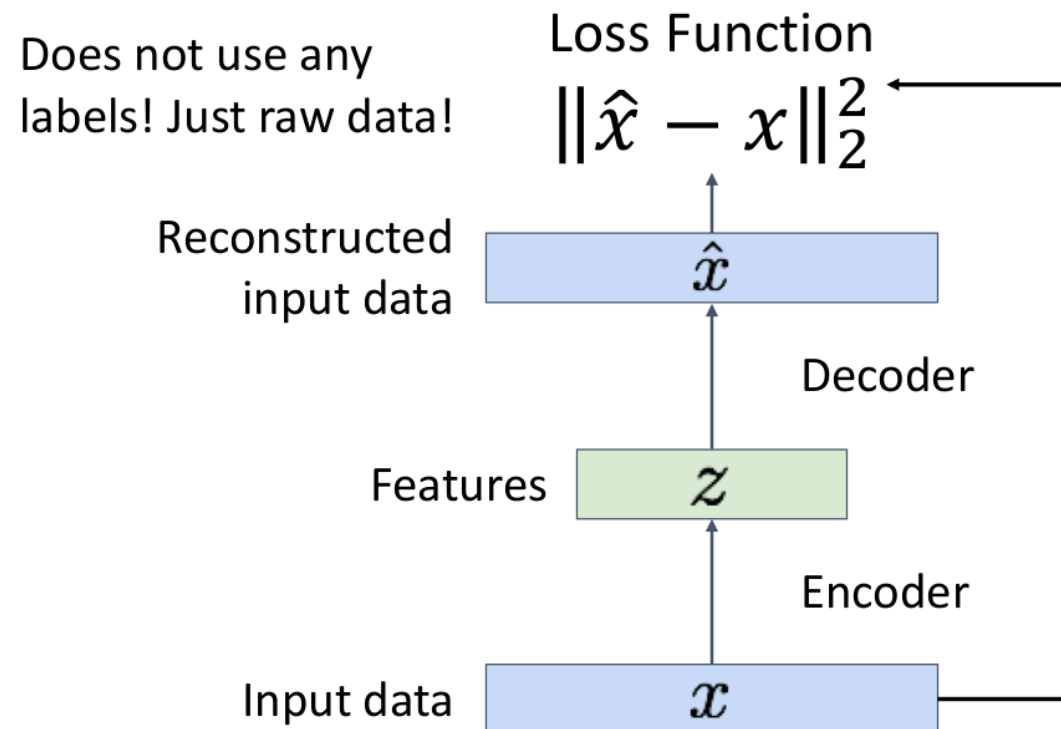
---

- In PixelRNN, PixelCNN, they defined parametric density function  $p(x) = f(x, W)$  and calculate for each input. And train the model to maximize this output.
- In VAE, Instead of maximizing the actual density value, **maximize the lower bound of the density.**
- $p_{\theta}(x) = \prod_{i=1}^n p_{\theta}(x_i | x_1, \dots, x_{i-1})$

# Variational Autoencoders (VAE)

## (Regular, non-variational) Autoencoders

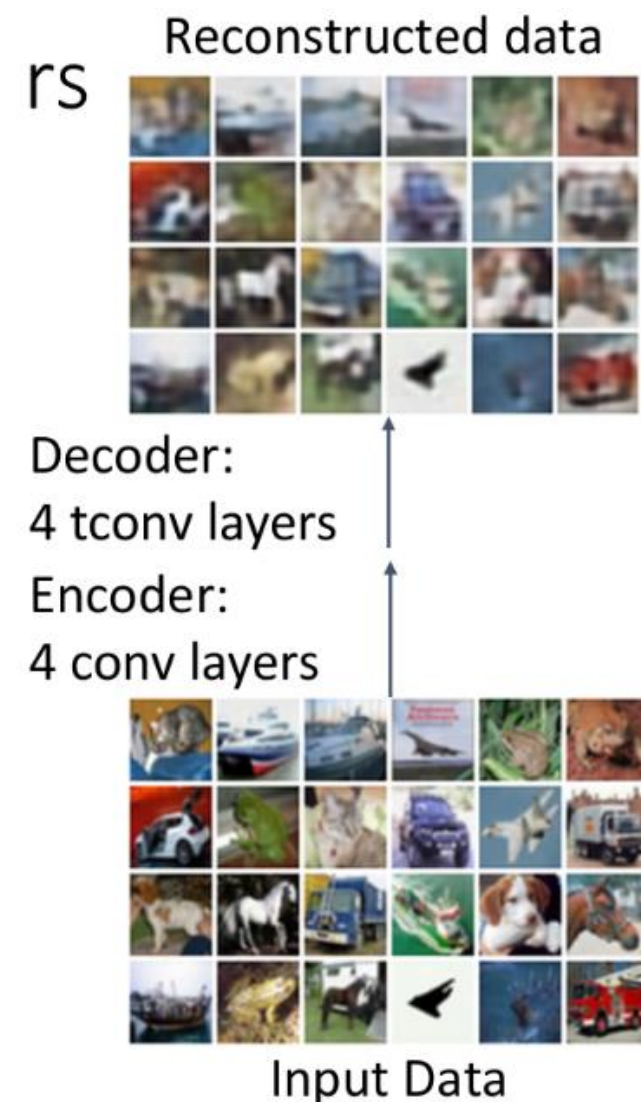
- Unsupervised method로, labels 없이 raw data  $x$ 로부터 feature vectors를 학습한다.  
(Unsupervised method for learning feature vectors from raw data  $x$ , without any labels)
- **Features** extracts useful information that can be used for downstream tasks.
- **Encoder** extracts features from input data,
- **Decoder** reconstruct the input data from the features.
- Encoder:
  - Originally: Linear + nonlinearity (sigmoid)
  - Later: Deep, fully-connected
  - Later: ReLU CNN (upconv)



# Variational Autoencoders (VAE)

## (Regular, non-variational) Autoencoders

- We expect the effect of **compressing input data through Encoder**.
- After learning, discard the decoder and use it for the downstream task using the encoder.
- **Not probabilistic** : 학습하지 않은 new data를 sampling할 수 없다.



# Variational Autoencoders (VAE)

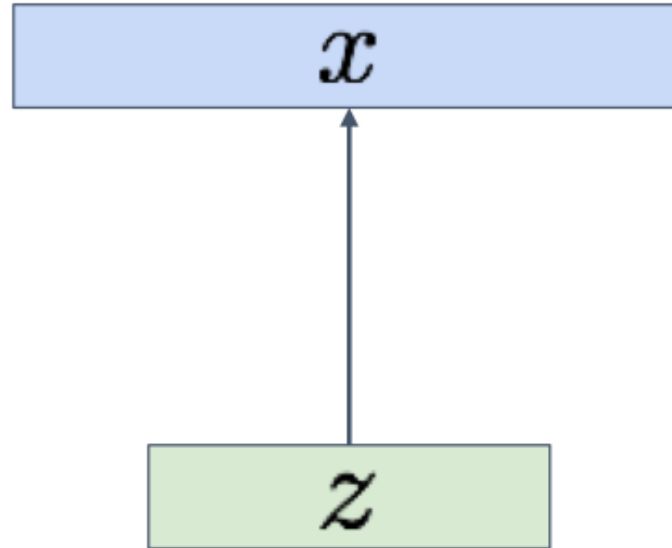
## Variational Autoencoders

Sample from  
conditional

$$p_{\theta^*}(x \mid z^{(i)})$$

Sample  $z$   
from prior

$$p_{\theta^*}(z)$$



- Autoencoder에 확률 개념을 도입하였다
- 1. raw data로부터 **latent features  $z$** 를 학습한다.
- 2. new data를 생성하기 위해 model로부터 **sampling**한다.

# Variational Autoencoders (VAE)

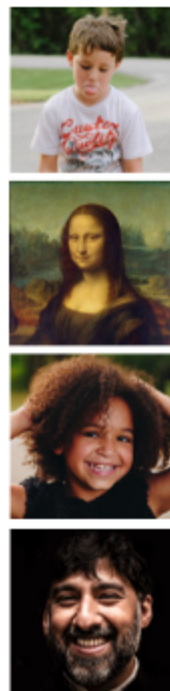
---

## Variational Autoencoders

- **Decoder:** Generating new data  $x$  from latent features  $z$  that is similar to input data but completely new.
- Sampling latent variables from prior distribution  $p_{\theta_*}(z)$ , put sampled  $z$  into the decoder to predict image  $x$
- At this time, output is not a single image but the distribution of images.
- $p_{\theta_*}(z)$  : prior distribution. PDF of  $x$ . (Gaussian distribution)
- $p_{\theta_*}(x|z^{(i)})$  : 주어진  $z$ 에서 특정  $x$ 가 나올 조건부 확률에 대한 PDF
- $\theta$  : parameter

# Variational Autoencoders (VAE)

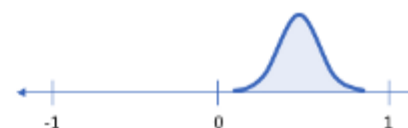
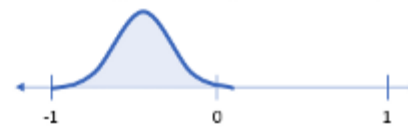
## Variational Autoencoders



Smile (discrete value)



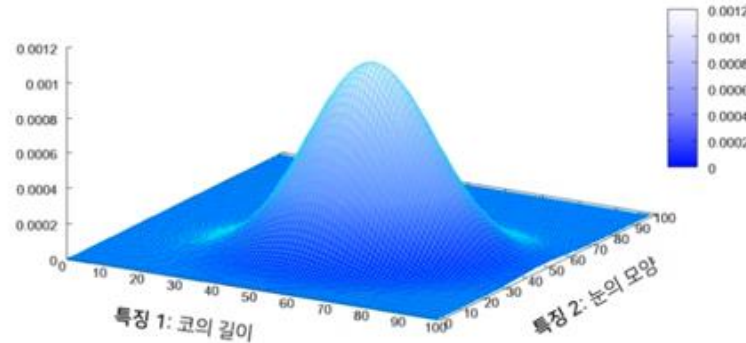
Smile (probability distribution)



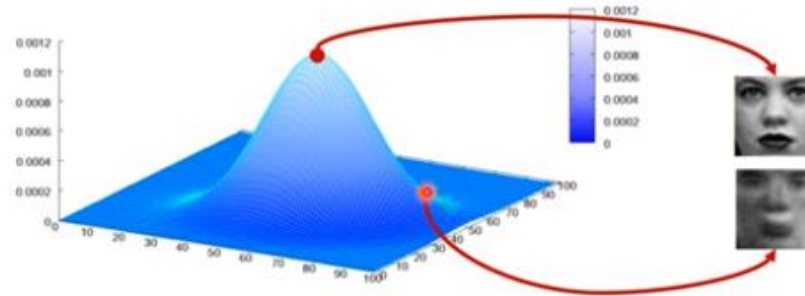
vs.

- VAE에서의  $z$ 는 AE에서의  $z$  (a value: low dimension of input data)와 다르게, 가우시안 확률분포에 기반한 확률값으로 나타낸다.
- input image가 들어오면 그 이미지의 다양한 특징들이 각각의 확률변수가 되는 확률분포를 만든다. 이 확률 분포를 잘 찾아내어 확률값이 높은 부분을 사용하면 그럴듯한 새로운 이미지를 생성할 수 있다.

# Variational Autoencoders (VAE)



입력 데이터의  
분포를 잘 근사하는 모델을 생성



- 이때 각 feature가 가우시안 분포를 따른다고 가정하고 latent  $z$ 는 각 feature의 평균과 분산값을 나타낸다.
- 예를 들어 한국인의 얼굴을 그리기 위해 눈, 코, 입 등의 feature를 Latent vector  $z$ 에 담고, 그  $z$ 를 이용해 그럴듯한 한국인의 얼굴을 그려내는 것이다. latent vector  $z$ 는 한국인 눈 모양의 평균 및 분산, 한국인 코 길이의 평균 및 분산, 한국인 머리카락 길이의 평균 및 분산 등등의 정보를 담고 있다고 생각할 수 있다.



# Generative Models II

Jisang Han

onground@korea.ac.kr

Artificial Intelligence in KU (AIKU)

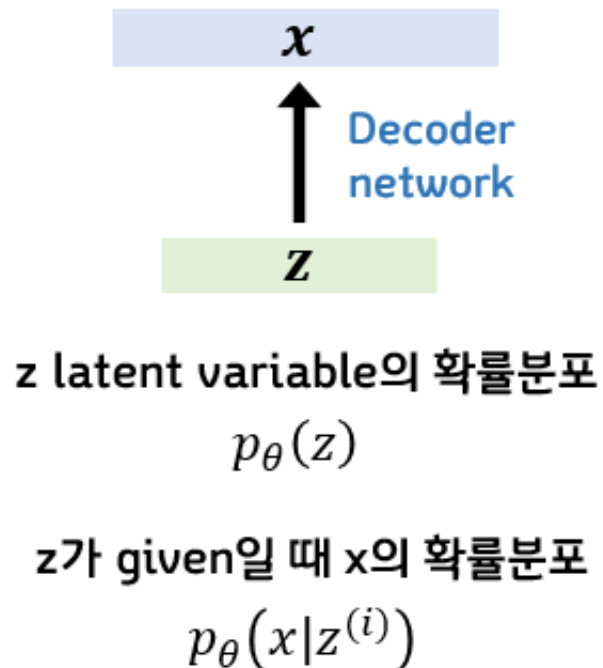
Department of Computer Science and Engineering, Korea University

**AIKU**

# Mathematical meaning

- If a model parameter  $\theta$  is given, the higher  $p_{\theta_*}(x)$  (The probability that the answer we want is  $x$ ) the better model.
- Train parameters to maximize  $p_{\theta_*}(x)$

## Decoder



## 어떻게 학습?

네트워크의 출력값이 있을 때  
우리가 원하는 정답  $x$ 가 나올 확률이 높길바람  
=  $x$ 의 likelihood를 최대화하는 확률분포 찾자

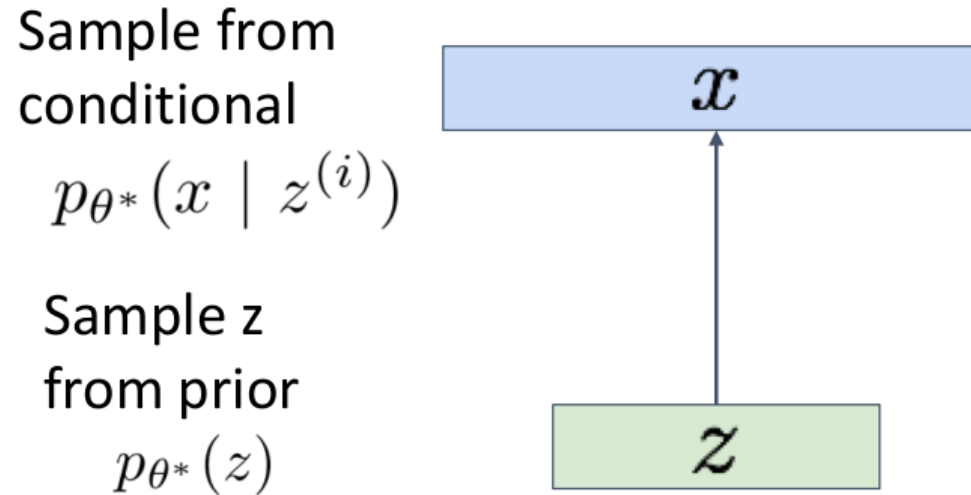


## Maximize

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)dz$$

# Variational Autoencoders

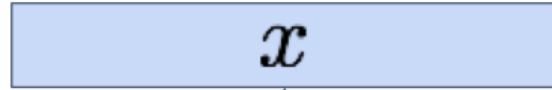
---



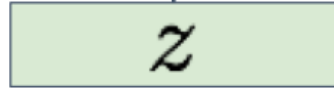
- When image comes in,  $p(z)$  is the appropriate Gaussian distribution for each pixel of image.
- That is, each pixel has a Gaussian distribution with  $\mu, \sigma$
- But if it's a **high resolution** image, it's going to have a lot of values. Therefore, the **diagonal gaussian distribution** is used instead of the general Gaussian distribution. That is, when  $z$  is given, there is no covariance between pixels of the generated image. **Pixels are independent.**

# Variational Autoencoders

Sample from  
conditional  
 $p_{\theta^*}(x \mid z^{(i)})$



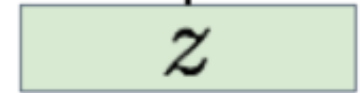
Sample  $z$   
from prior  
 $p_{\theta^*}(z)$



Sample from  
conditional  
 $p_{\theta^*}(x \mid z^{(i)})$



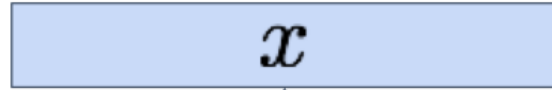
Sample  $z$   
from prior  
 $p_{\theta^*}(z)$



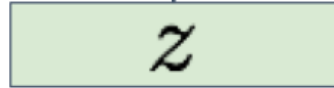
- **Decoder** outputs mean  $\mu_{x|z}$  and diagonal covariance  $\Sigma_{x|z}$  for the input  $z$
- Then sample  $x$  from the above Gaussian distribution

# Variational Autoencoders - Train

Sample from  
conditional  
 $p_{\theta^*}(x | z^{(i)})$



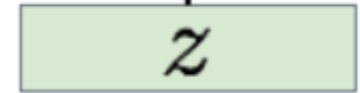
Sample  $z$   
from prior  
 $p_{\theta^*}(z)$



Sample from  
conditional  
 $p_{\theta^*}(x | z^{(i)})$



Sample  $z$   
from prior  
 $p_{\theta^*}(z)$



- Maximize likelihood of data  $p_{\theta_*}(x)$

$$p_{\theta}(x) = \frac{p_{\theta}(x|z)p_{\theta}(z)}{p_{\theta}(z|x)} \quad q_{\phi}(z|x) \approx p_{\theta}(z|x) \quad p_{\theta}(x) = \frac{p_{\theta}(x|z)p_{\theta}(z)}{p_{\theta}(z|x)} \approx \frac{p_{\theta}(x|z)p_{\theta}(z)}{q_{\phi}(z|x)}$$

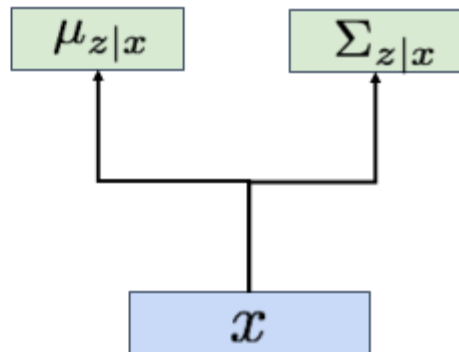
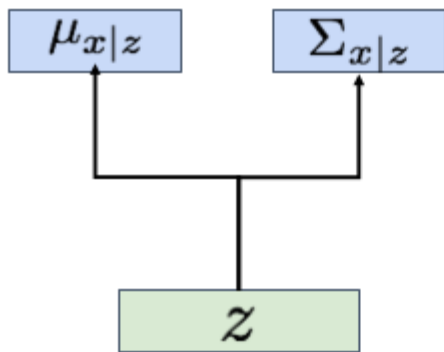
# Variational Autoencoders - Train

**Decoder network** inputs  
latent code  $z$ , gives  
distribution over data  $x$

**Encoder network** inputs  
data  $x$ , gives distribution  
over latent codes  $z$

$$p_{\theta}(x | z) = N(\mu_{x|z}, \Sigma_{x|z})$$

$$q_{\phi}(z | x) = N(\mu_{z|x}, \Sigma_{z|x})$$



$$\log p_{\theta}(x) = \log \frac{p_{\theta}(x|z)p(z)}{p_{\theta}(z|x)}$$

$$= \log \frac{p_{\theta}(x|z)p(z)q_{\phi}(z|x)}{p_{\theta}(z|x)q_{\phi}(z|x)}$$

$$= \log p_{\theta}(x|z) - \log \frac{q_{\phi}(z|x)}{p(z)} + \log \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)}$$

wrap in an expectation since it doesn't depend on  $z$

$$\log p_{\theta}(x) = E_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x)]$$

$$= E_z[\log p_{\theta}(x|z)] - E_z \left[ \log \frac{q_{\phi}(z|x)}{p(z)} \right] + E_z \left[ \log \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} \right]$$

$$= E_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{KL}(q_{\phi}(z|x), p(z)) + D_{KL}(q_{\phi}(z|x), p_{\theta}(z|x))$$

# Variational Autoencoders - Train

---

$$\log p_{\theta}(x) = E_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x)]$$

$$= E_z[\log p_{\theta}(x|z)] - E_z \left[ \log \frac{q_{\phi}(z|x)}{p(z)} \right] + E_z \left[ \log \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} \right]$$

$$= E_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{KL}(q_{\phi}(z|x), p(z)) + D_{KL}(q_{\phi}(z|x), p_{\theta}(z|x))$$

$$\log p_{\theta}(x) \geq E_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{KL}(q_{\phi}(z|x), p(z))$$

Lower bound of likelihood

Through this, encoders and decoders are learned jointly to maximize the variable lower bound of data-like hood.



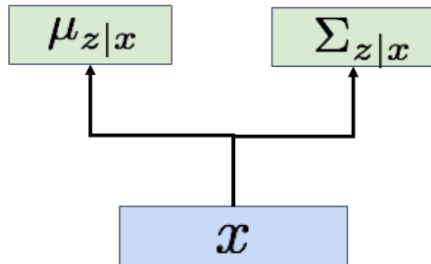
# Variational Autoencoders

Jointly train **encoder**  $q$  and **decoder**  $p$  to maximize the **variational lower bound** on the data likelihood

$$\log p_{\theta}(x) \geq E_{z \sim q_{\phi}(z|x)} [\log p_{\theta}(x|z)] - D_{KL} \left( q_{\phi}(z|x), p(z) \right)$$

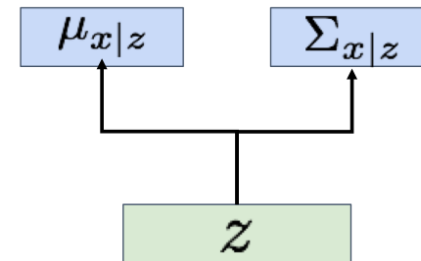
Encoder Network

$$q_{\phi}(z | x) = N(\mu_{z|x}, \Sigma_{z|x})$$



Decoder Network

$$p_{\theta}(x | z) = N(\mu_{x|z}, \Sigma_{x|z})$$



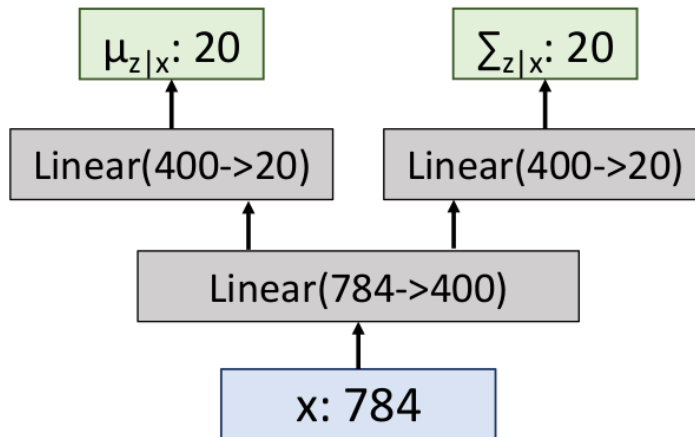
# Example : Fully connected VAE

MNIST Dataset

- $x = 28 \times 28$  image, flattened to 784-dim vector
- $z = 20$ -dim vector (hyper-parameter)

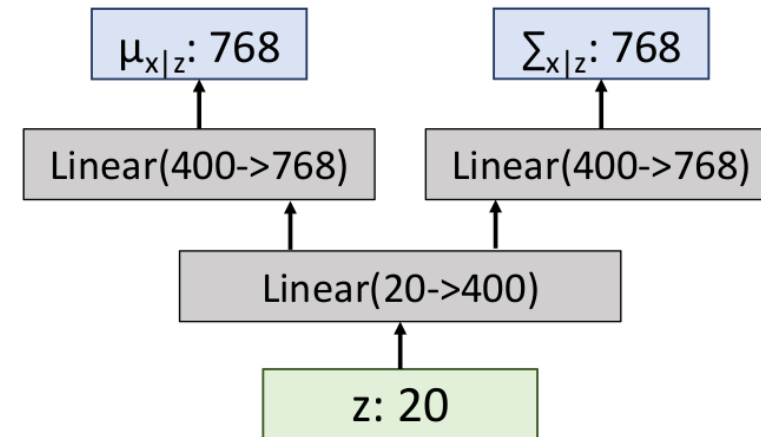
**Encoder Network**

$$q_{\phi}(z | x) = N(\mu_{z|x}, \Sigma_{z|x})$$



**Decoder Network**

$$p_{\theta}(x | z) = N(\mu_{x|z}, \Sigma_{x|z})$$



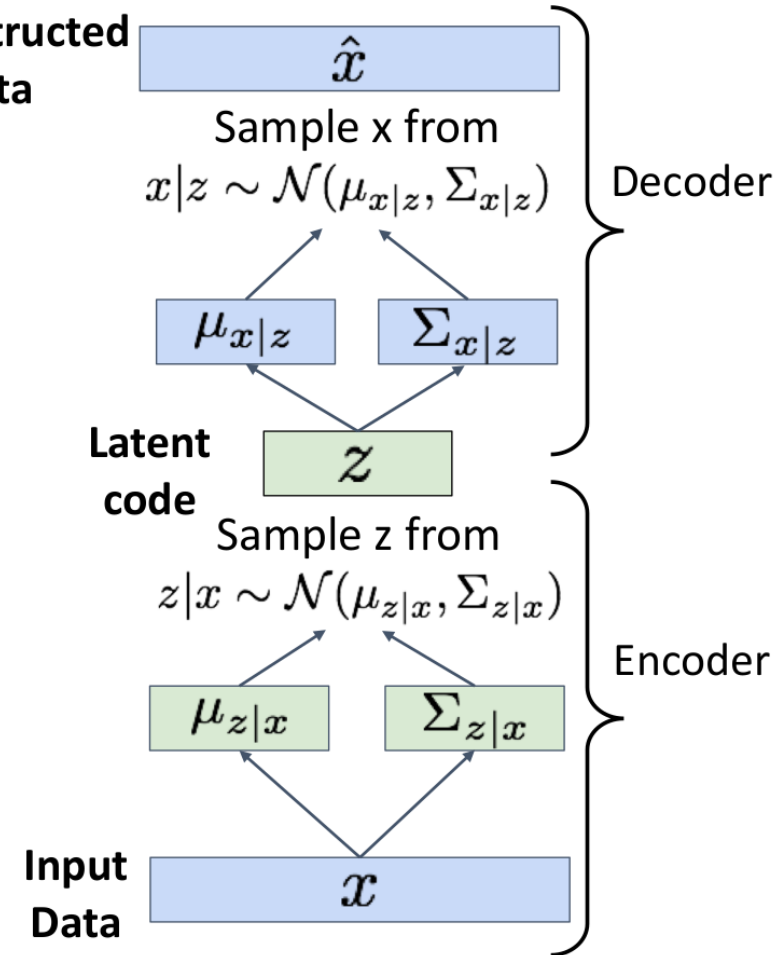
# Example : Fully connected VAE

## Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{KL}(q_{\phi}(z|x), p(z))$$

1. Run input data through **encoder** to get a distribution over latent codes
2. **Encoder output should match the prior  $p(z)$ !**
3. Sample code  $z$  from encoder output
4. Run sampled code through **decoder** to get a distribution over data samples
5. **Original input data should be likely under the distribution output from (4)!**
6. Can sample a reconstruction from (4)



# Variational Autoencoders

32x32 CIFAR-10



Labeled Faces in the Wild



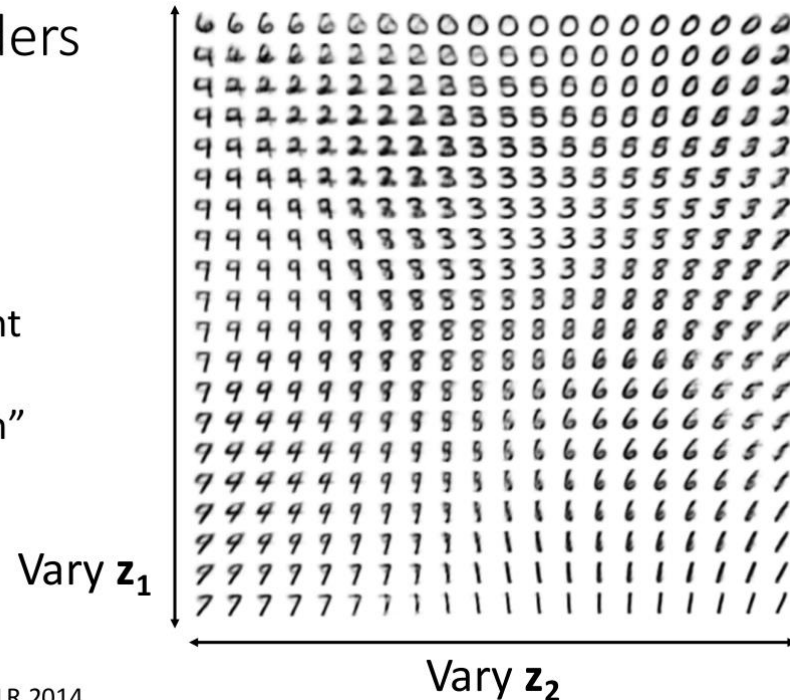
# Variational Autoencoders

- VAE : Editing with  $z$

Variational Autoencoders

The diagonal prior on  $p(z)$  causes  
dimensions of  $z$  to be independent

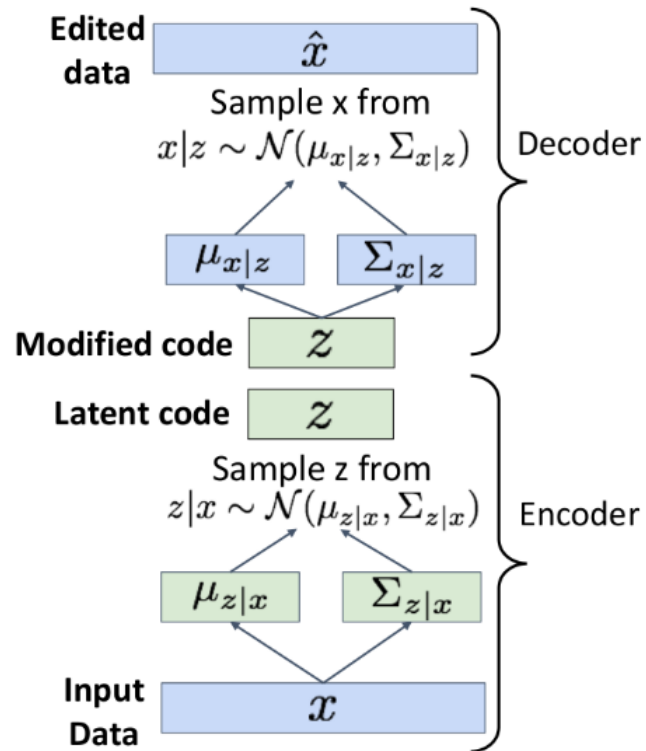
“Disentangling factors of variation”



Kingma and Welling, Auto-Encoding Variational Bayes, ICLR 2014

# Variational Autoencoders

- VAE : Editing with  $z$





# Variational Autoencoders

- VAE : Editing with  $z$

## Variational Autoencoders

The diagonal prior on  $p(z)$  causes dimensions of  $z$  to be independent

“Disentangling factors of variation”

Kingma and Welling, Auto-Encoding Variational Bayes, ICLR 2014

Degree of smile

Vary  $z_1$

Head pose

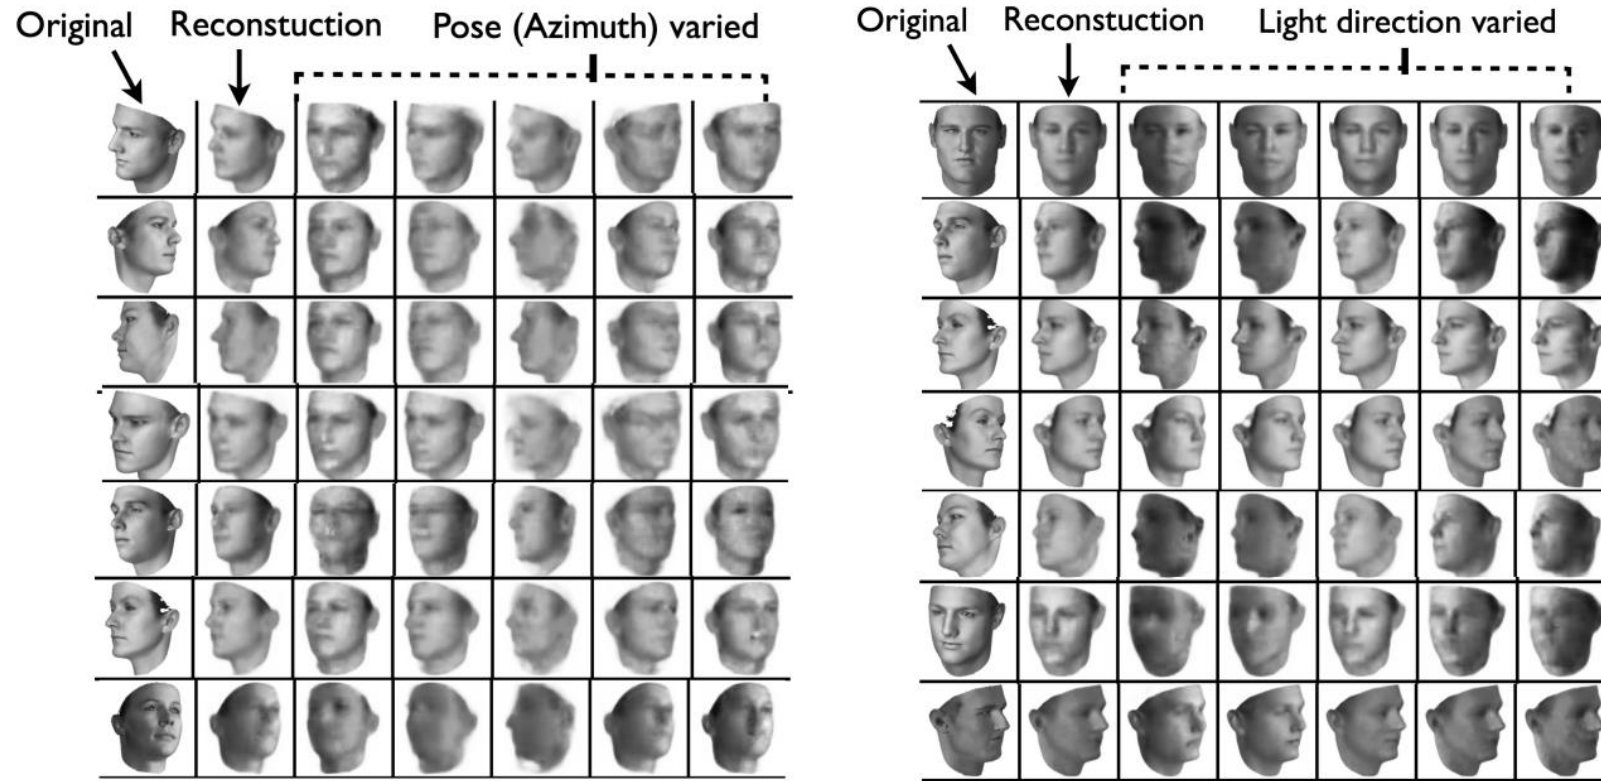
Vary  $z_2$





# Variational Autoencoders

- VAE : Editing with  $z$



# Variational Autoencoders

---

- **VAE : Summary**

Probabilistic spin to traditional autoencoders => allows generating data

Defines an intractable density => derive and optimize a (variational) lower bound

**Pros:**

- Principled approach to generative models
- Allows inference of  $q(z|x)$ , can be useful feature representation for other tasks

**Cons:**

- Maximizes lower bound of likelihood: okay, but not as good evaluation as PixelRNN/PixelCNN
- Samples blurrier and lower quality compared to state-of-the-art (GANs)

**Active areas of research:**

- More flexible approximations, e.g. richer approximate posterior instead of diagonal Gaussian, e.g., Gaussian Mixture Models (GMMs)
- Incorporating structure in latent variables, e.g., Categorical Distributions

**Autoregressive models**

- Directly maximize  $p(\text{data})$
- High-quality generated images
- Slow to generate images
- No explicit latent codes

**Variational models**

- Maximize lower-bound on  $p(\text{data})$
- Generated images often blurry
- Very fast to generate images
- Learn rich latent codes

# Generative Adversarial Networks

Models	설명
Autoregressive models	training data의 likelihood를 직접(directly) maximize한다. $p_{\theta}(x) = \prod_{i=1}^N p_{\theta}(x_i   x_1, \dots, x_{i-1})$
VAE	latent $z$ 를 추가하였고, likelihood의 lower bound를 maximize한다. $p_{\theta}(x) = \int_Z p_{\theta}(x z)p(z) dz \geq E_{z \sim q_{\phi}(z x)}[\log p_{\theta}(x z)] - D_{KL}(q_{\phi}(z x), p(z))$
GANs	$p(x)$ 를 모델링하는 것을 포기한다. 그러나 $p(x)$ 로부터 샘플링할 수 있도록 한다.

# Generative Adversarial Networks

---

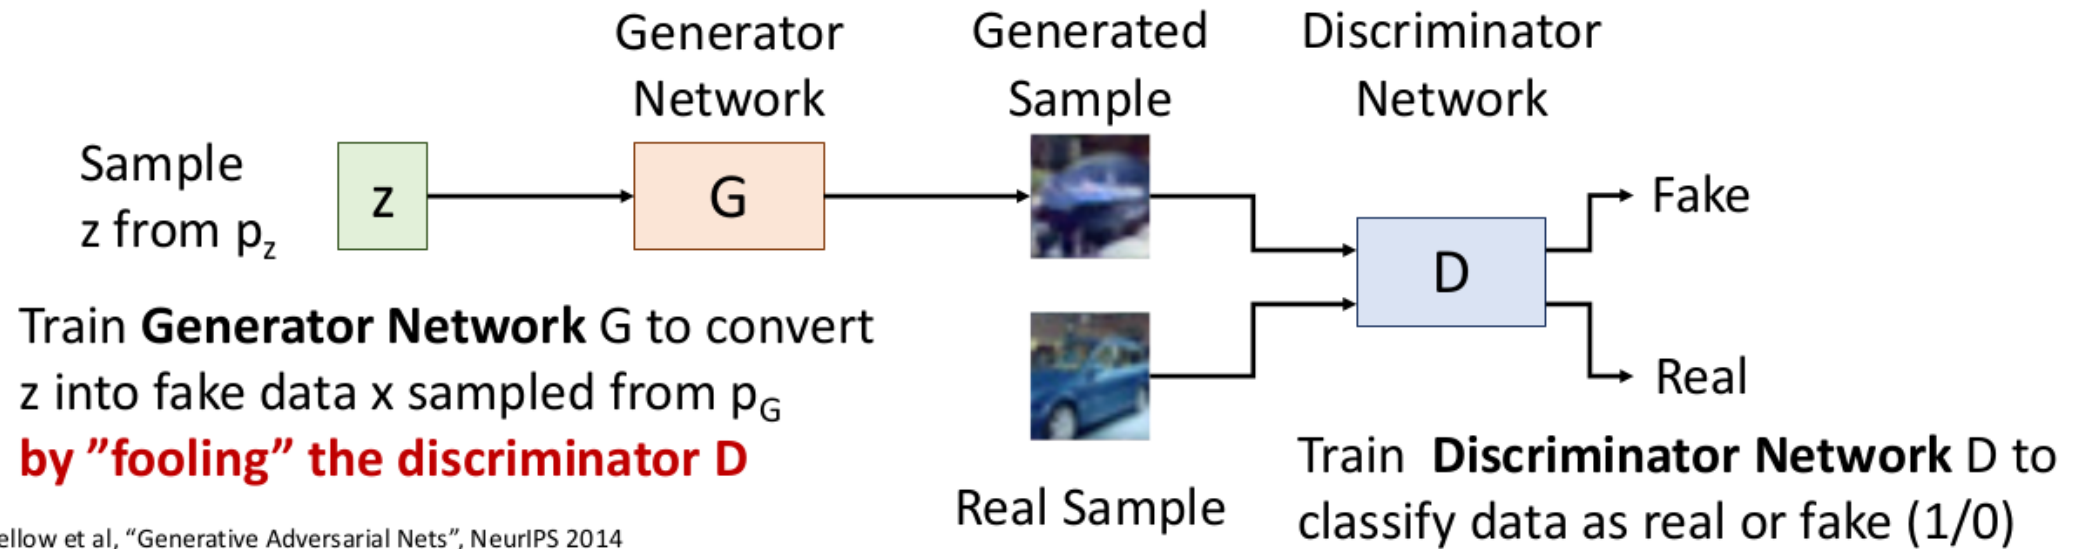
- **Setup**

- $p_{\text{data}}(x)$  : real data distribution
- $x_i$  : our train data from  $p_{\text{data}}(x)$

- **Idea**

- Suppose a latent variable  $z$  with  $p(z)$  which is a simple prior (diagonal Gaussian, uniform distribution, etc.).
- Sample  $z$  from  $p(z)$  and pass through **Generative Network  $G$** .
- $x = G(z)$
- Then the  $x$  is from Generative distribution  $p_G$ .
- Therefore we want  $p_G = p_{\text{data}}$  (Our generative distribution to be real data distribution)

# Generative Adversarial Networks



- **Generator**
  - By sampling  $x$  from  $p_G$ , train the model to generate an image that the Discriminator get fooled to think that the image was from  $p_{data}$ .
- **Discriminator**
  - Train to discriminate the generated sample and a real sample (real/fake(1/0)).
- Jointly train the two networks. Then  $p_G$  will converge to  $p_{data}$ .

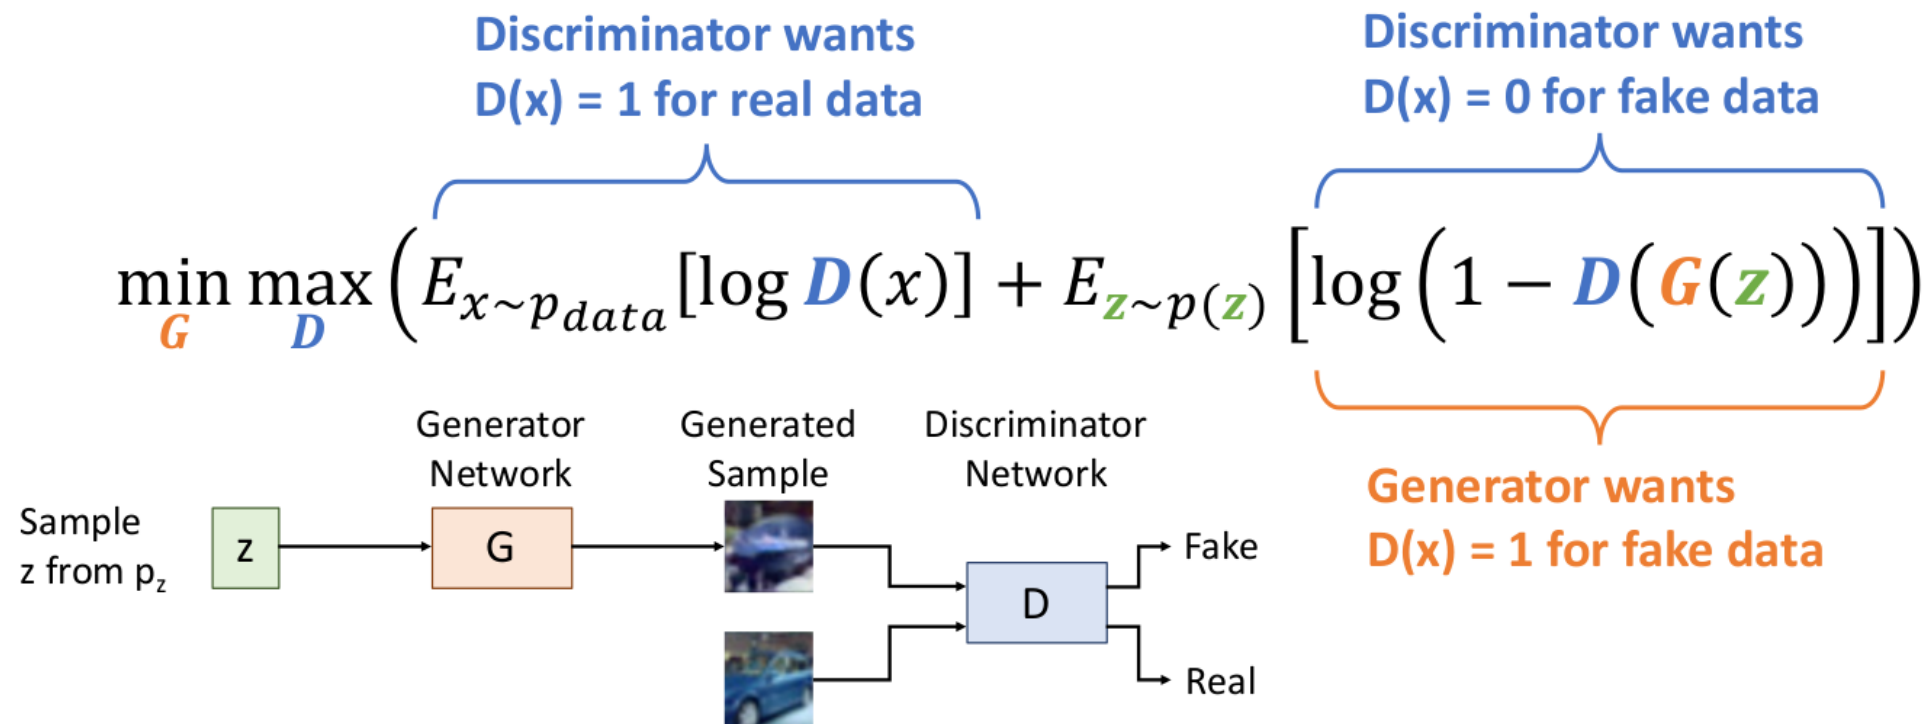
# GANs : Training Objective

---

- Train **Generator G** and **Discriminator D** jointly by **minmax game**.

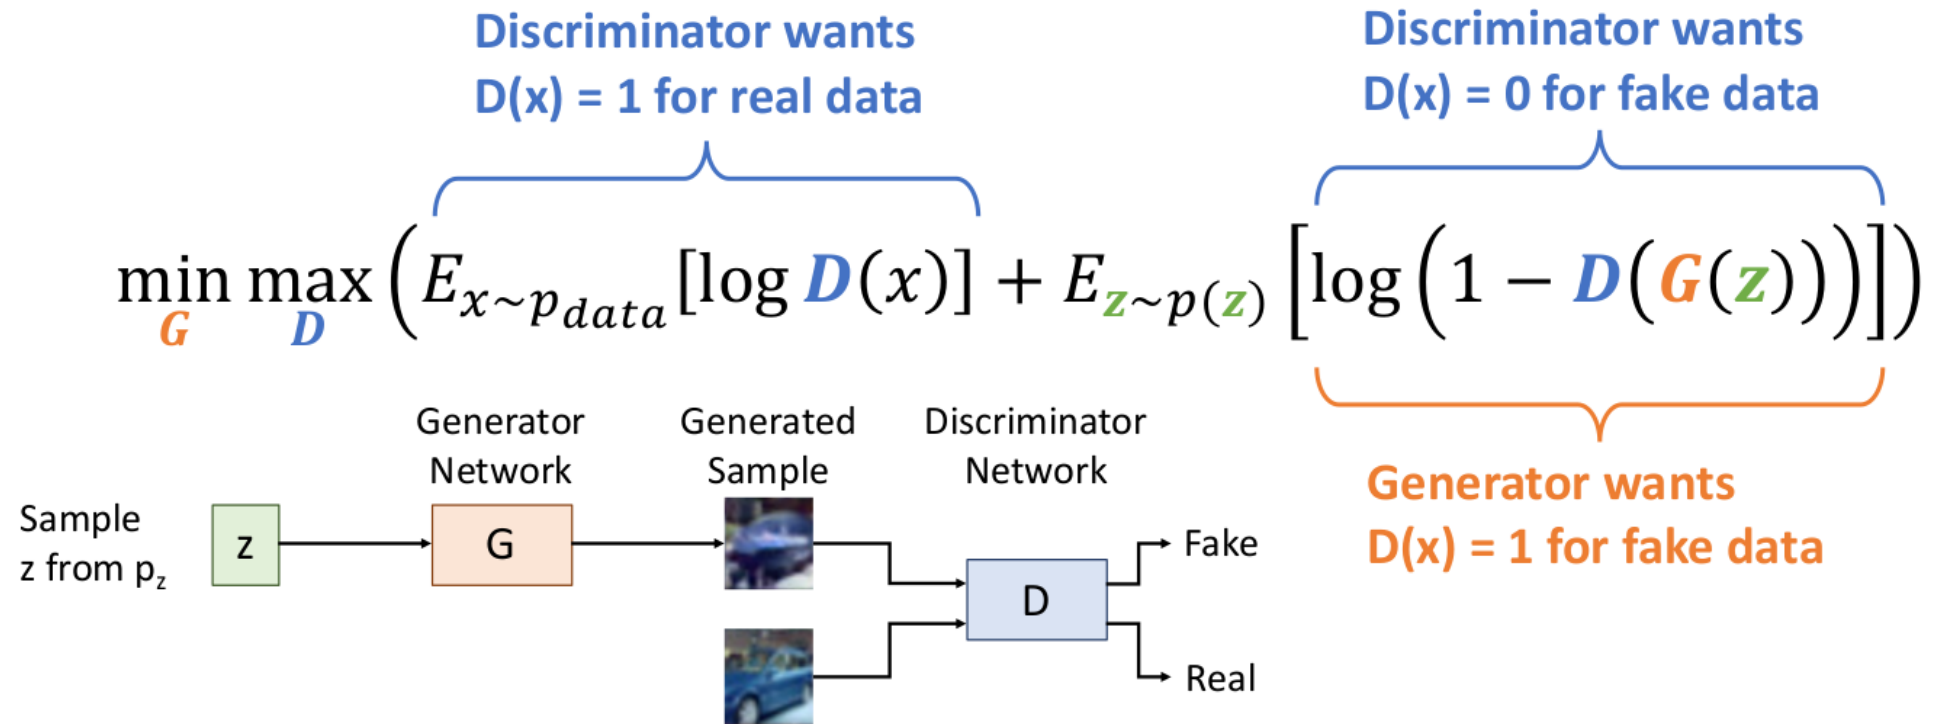
$$\min_G \max_D (E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))])$$

# GANs : Training Objective



- $x$  sampled from  $p_{data}$ , which is the real data to be REAL
- If  $D(x) < 1$ , it passes log term and becomes very small negative value. So we train  $D(x) = 1$  that the whole term can be maximized by  $D$ .

# GANs : Training Objective

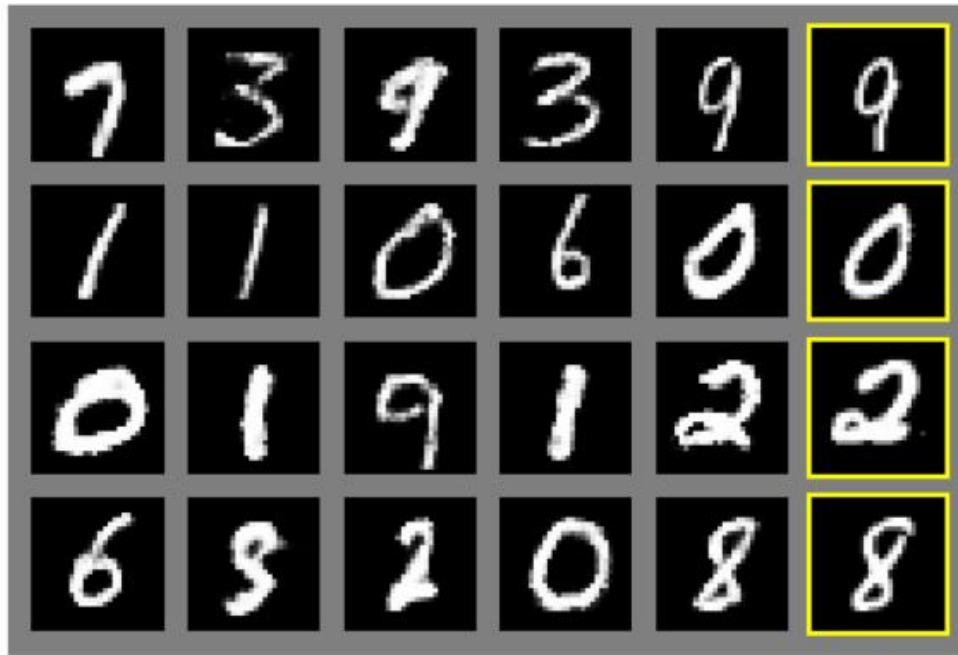


- $z$  sampled from  $p(z)$ , pass it to Generator  $G$ , and  $G$  outputs generated sample  $G(z)$ . Train Discriminator to discriminate the generated sample  $G(z)$  is fake. (fake to be FAKE)
- Generator  $G$  trains Discriminator  $D$  to discriminate  $G(z)$  is REAL. (fake to be REAL)



# GANs

Generated samples

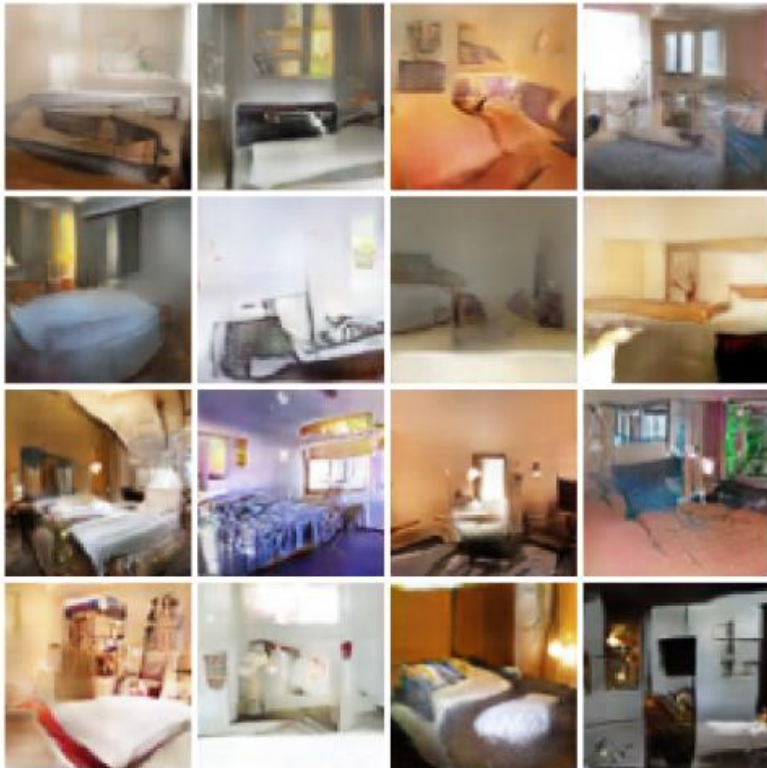


Nearest neighbor from training set

Goodfellow et al, "Generative Adversarial Nets", NeurIPS 2014

# GANs

## Wasserstein GAN (WGAN)



Arjovsky, Chintala, and Bottou, "Wasserstein GAN", 2017

## WGAN with Gradient Penalty (WGAN-GP)



Gulrajani et al, "Improved Training of Wasserstein GANs", NeurIPS 2017



# GANs

256 x 256 bedrooms



1024 x 1024 faces



Karras et al, "Progressive Growing of GANs for Improved Quality, Stability, and Variation", ICLR 2018



# GANs

512 x 384 cars



1024 x 1024 faces

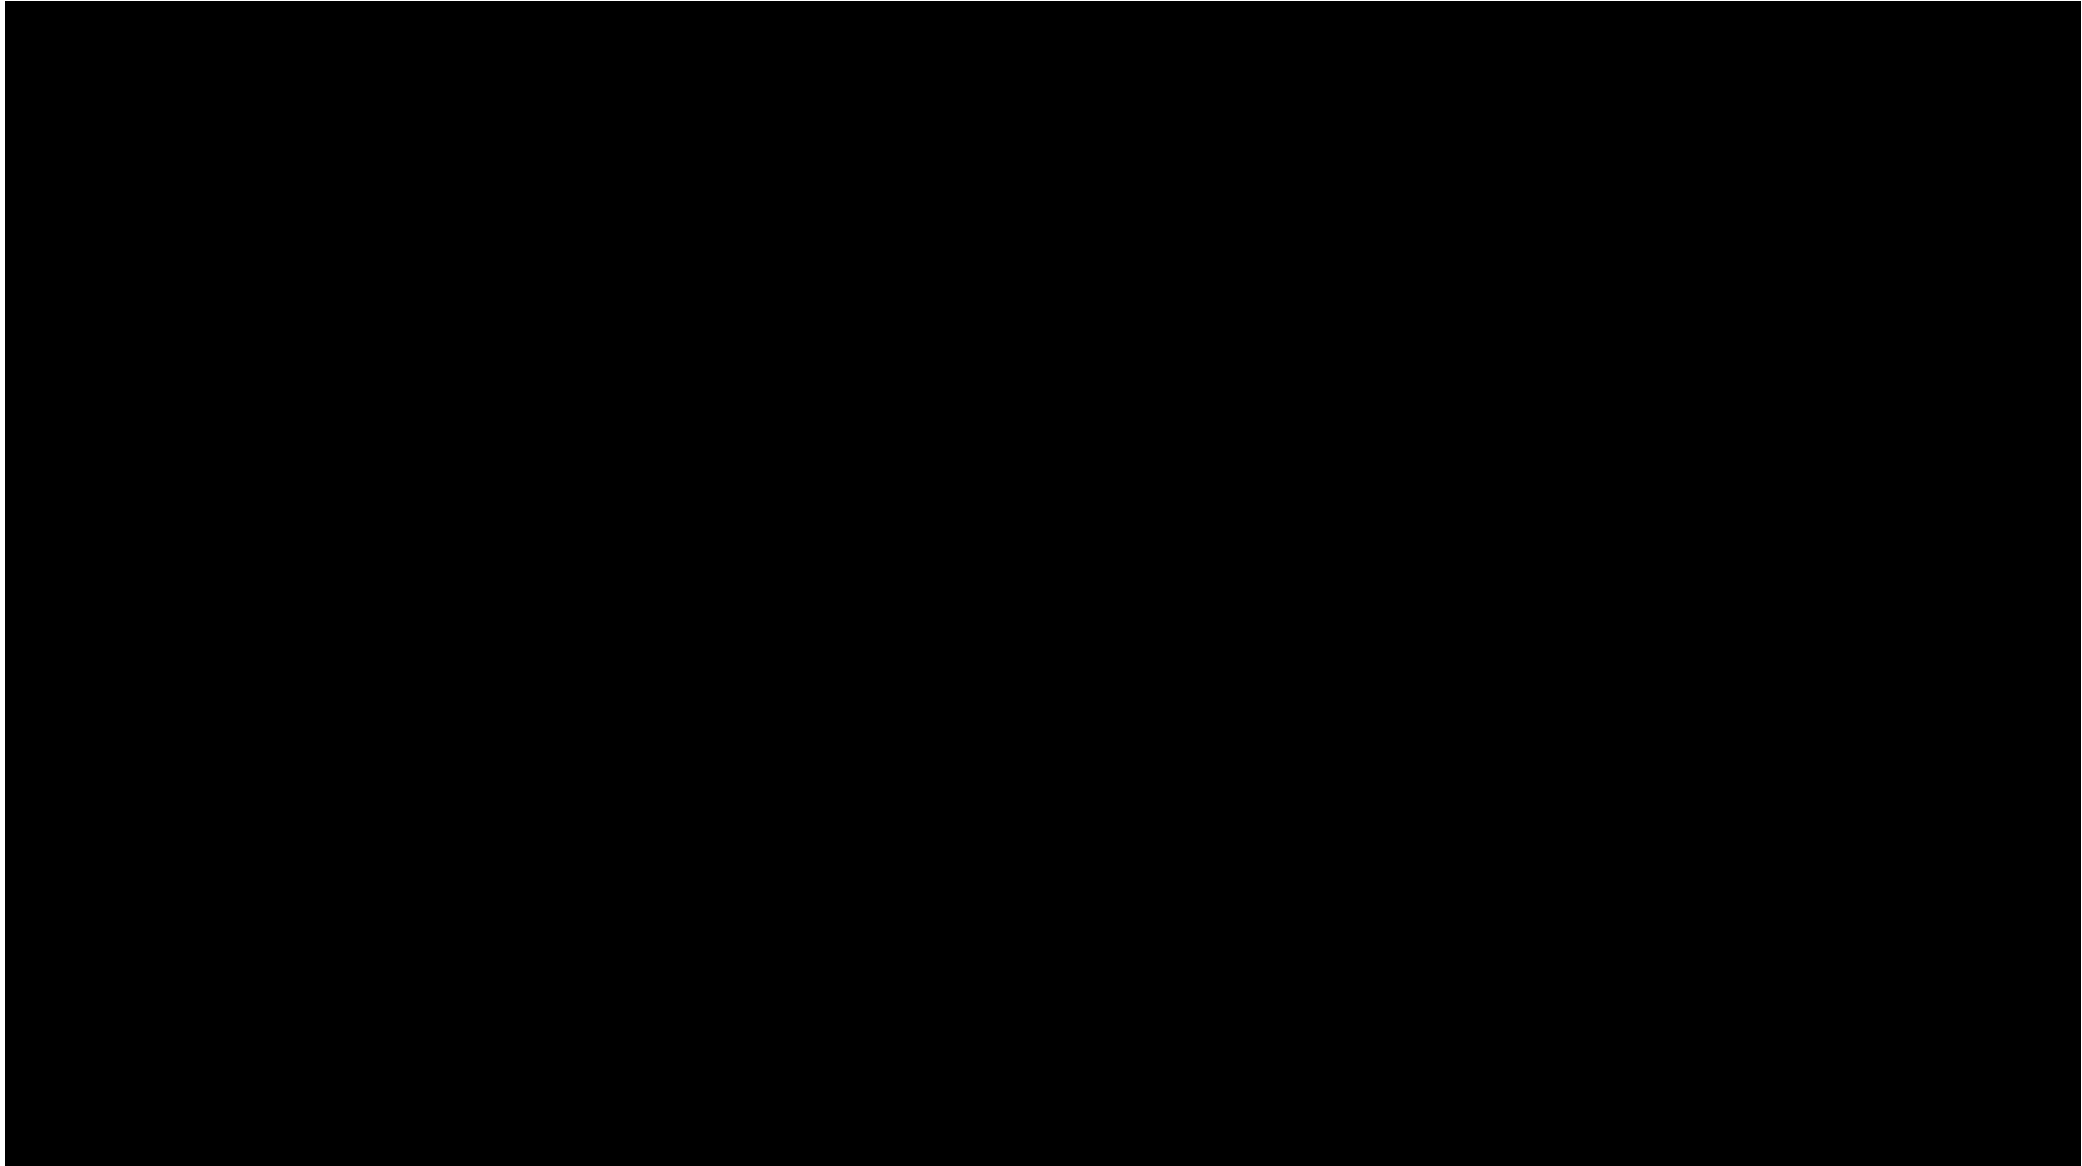


Karras et al, "A Style-Based Generator Architecture for Generative Adversarial Networks", CVPR 2019

[Images](#) are licensed under [CC BY-NC 4.0](#)

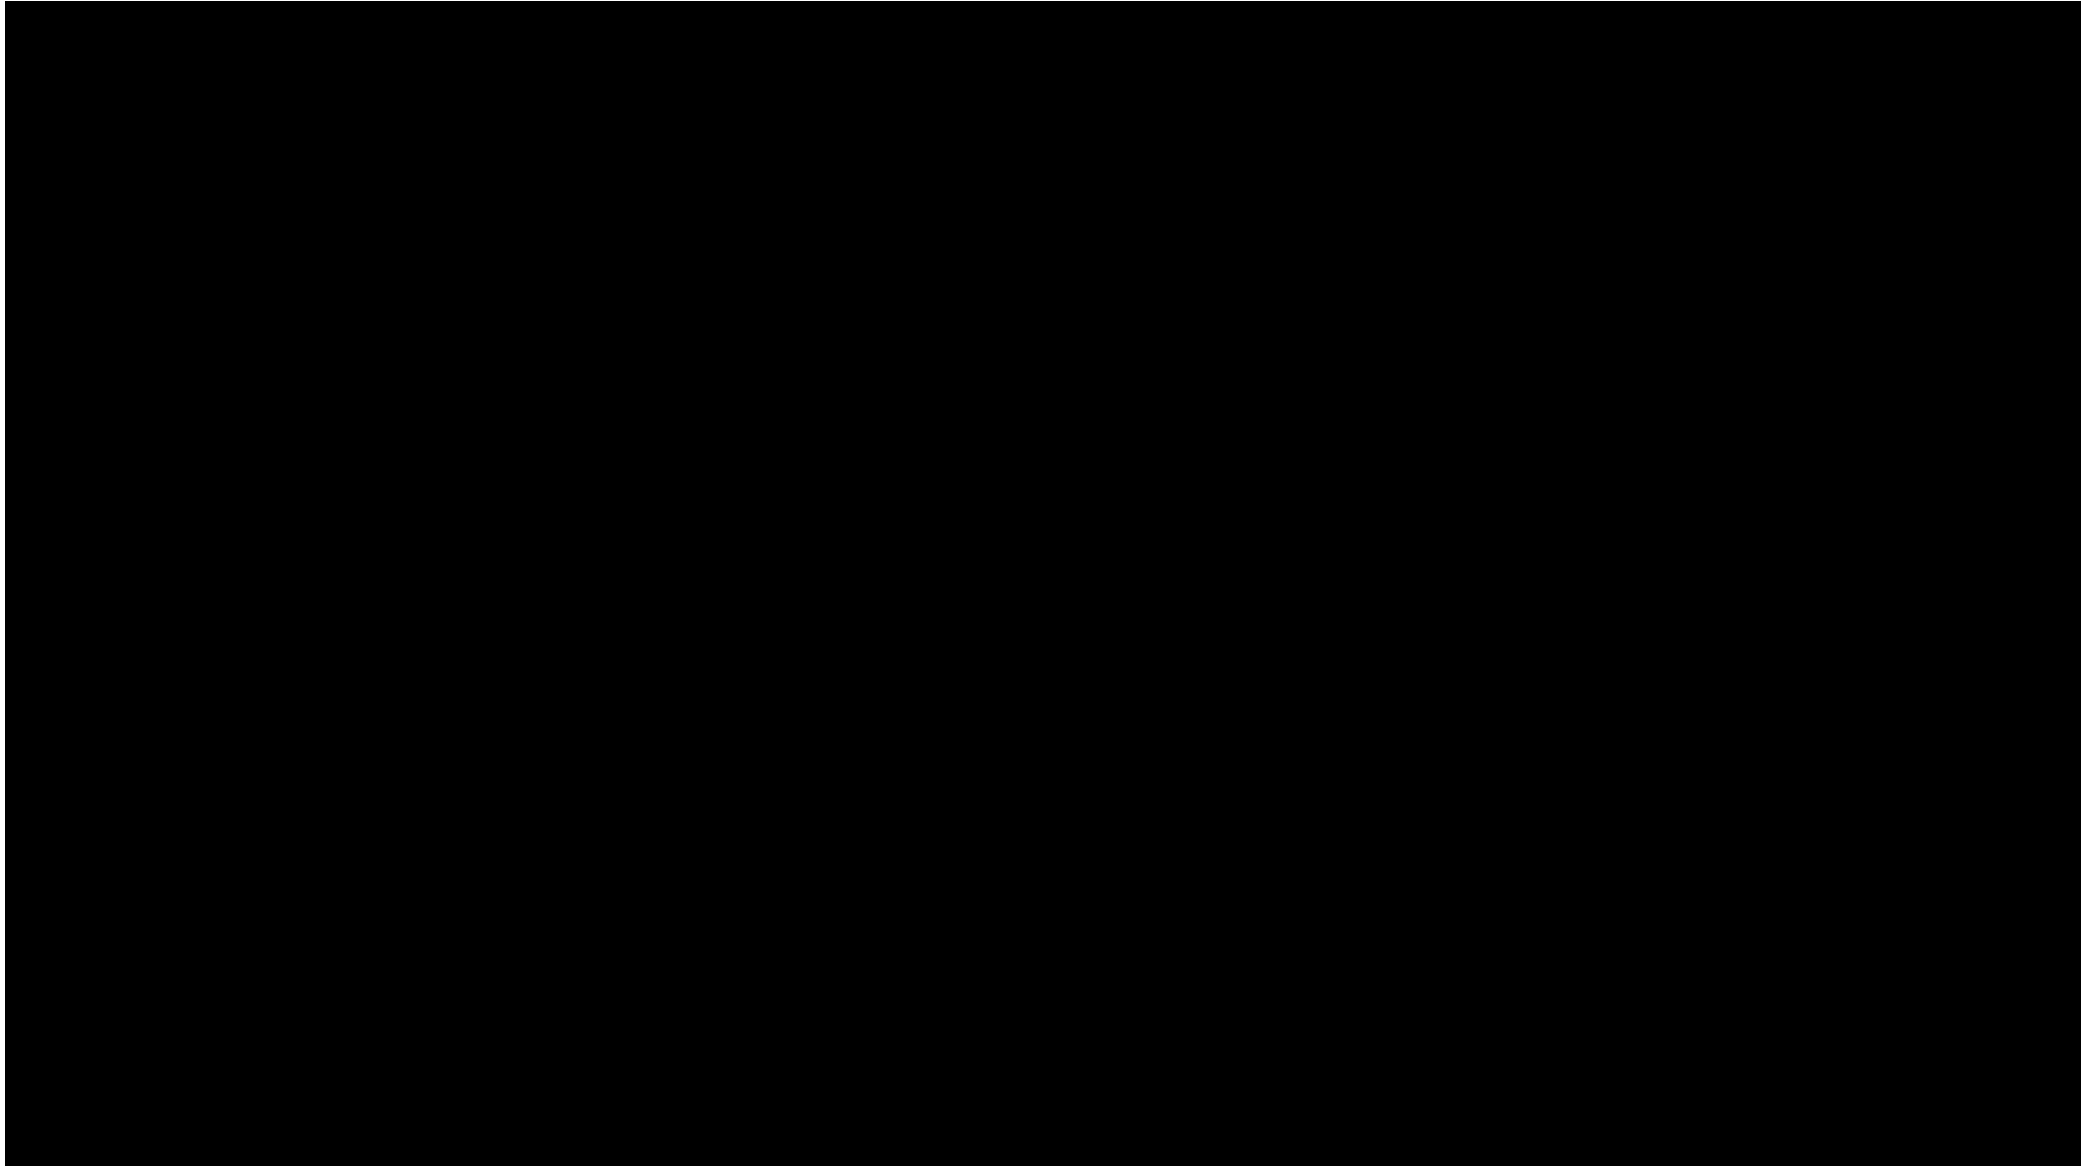
# GANs

---



# GANs

---



# GANs

This bird is red and brown in color, with a stubby beak



The bird is short and stubby with yellow on its body



A bird with a medium orange bill white body and gray wings and webbed feet



This small black bird has a short, slightly curved bill and long legs



A picture of a very clean living room



A group of people on skis stand in the snow



Eggs fruit candy nuts and meat served on white dish



A street sign on a stoplight pole in the middle of a day



Zhang et al, "StackGAN++: Realistic Image Synthesis with Stacked Generative Adversarial Networks.", TPAMI 2018

Zhang et al, "StackGAN: Text to Photo-realistic Image Synthesis with Stacked Generative Adversarial Networks.", ICCV 2017

Reed et al, "Generative Adversarial Text-to-Image Synthesis", ICML 2016



# GANs

bicubic  
(21.59dB/0.6423)



SRResNet  
(23.53dB/0.7832)



SRGAN  
(21.15dB/0.6868)



original



Ledig et al, "Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network", CVPR 2017

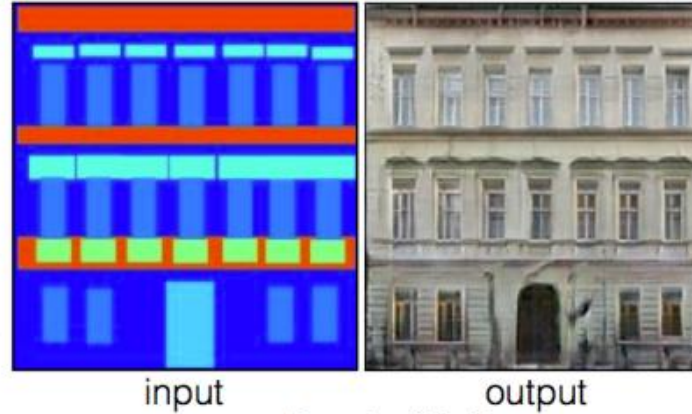


# GANs

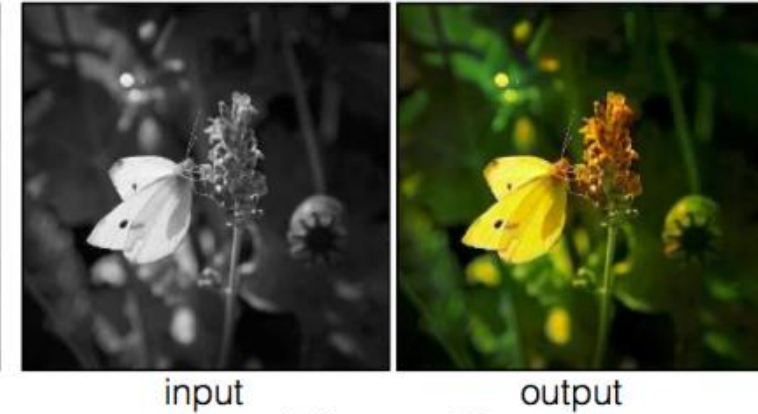
Labels to Street Scene



Labels to Facade



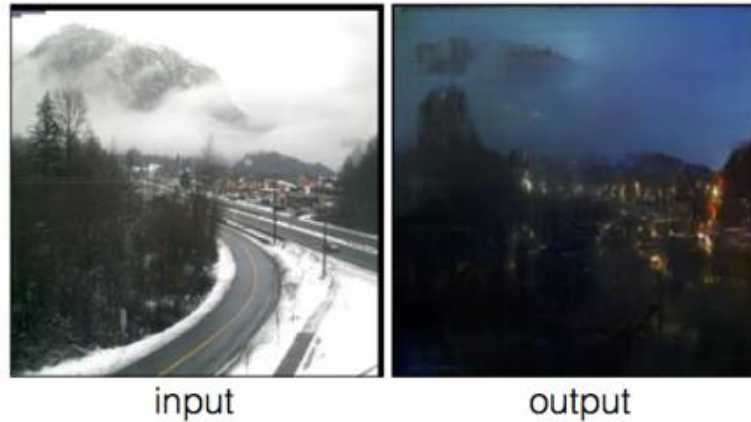
BW to Color



Aerial to Map



Day to Night



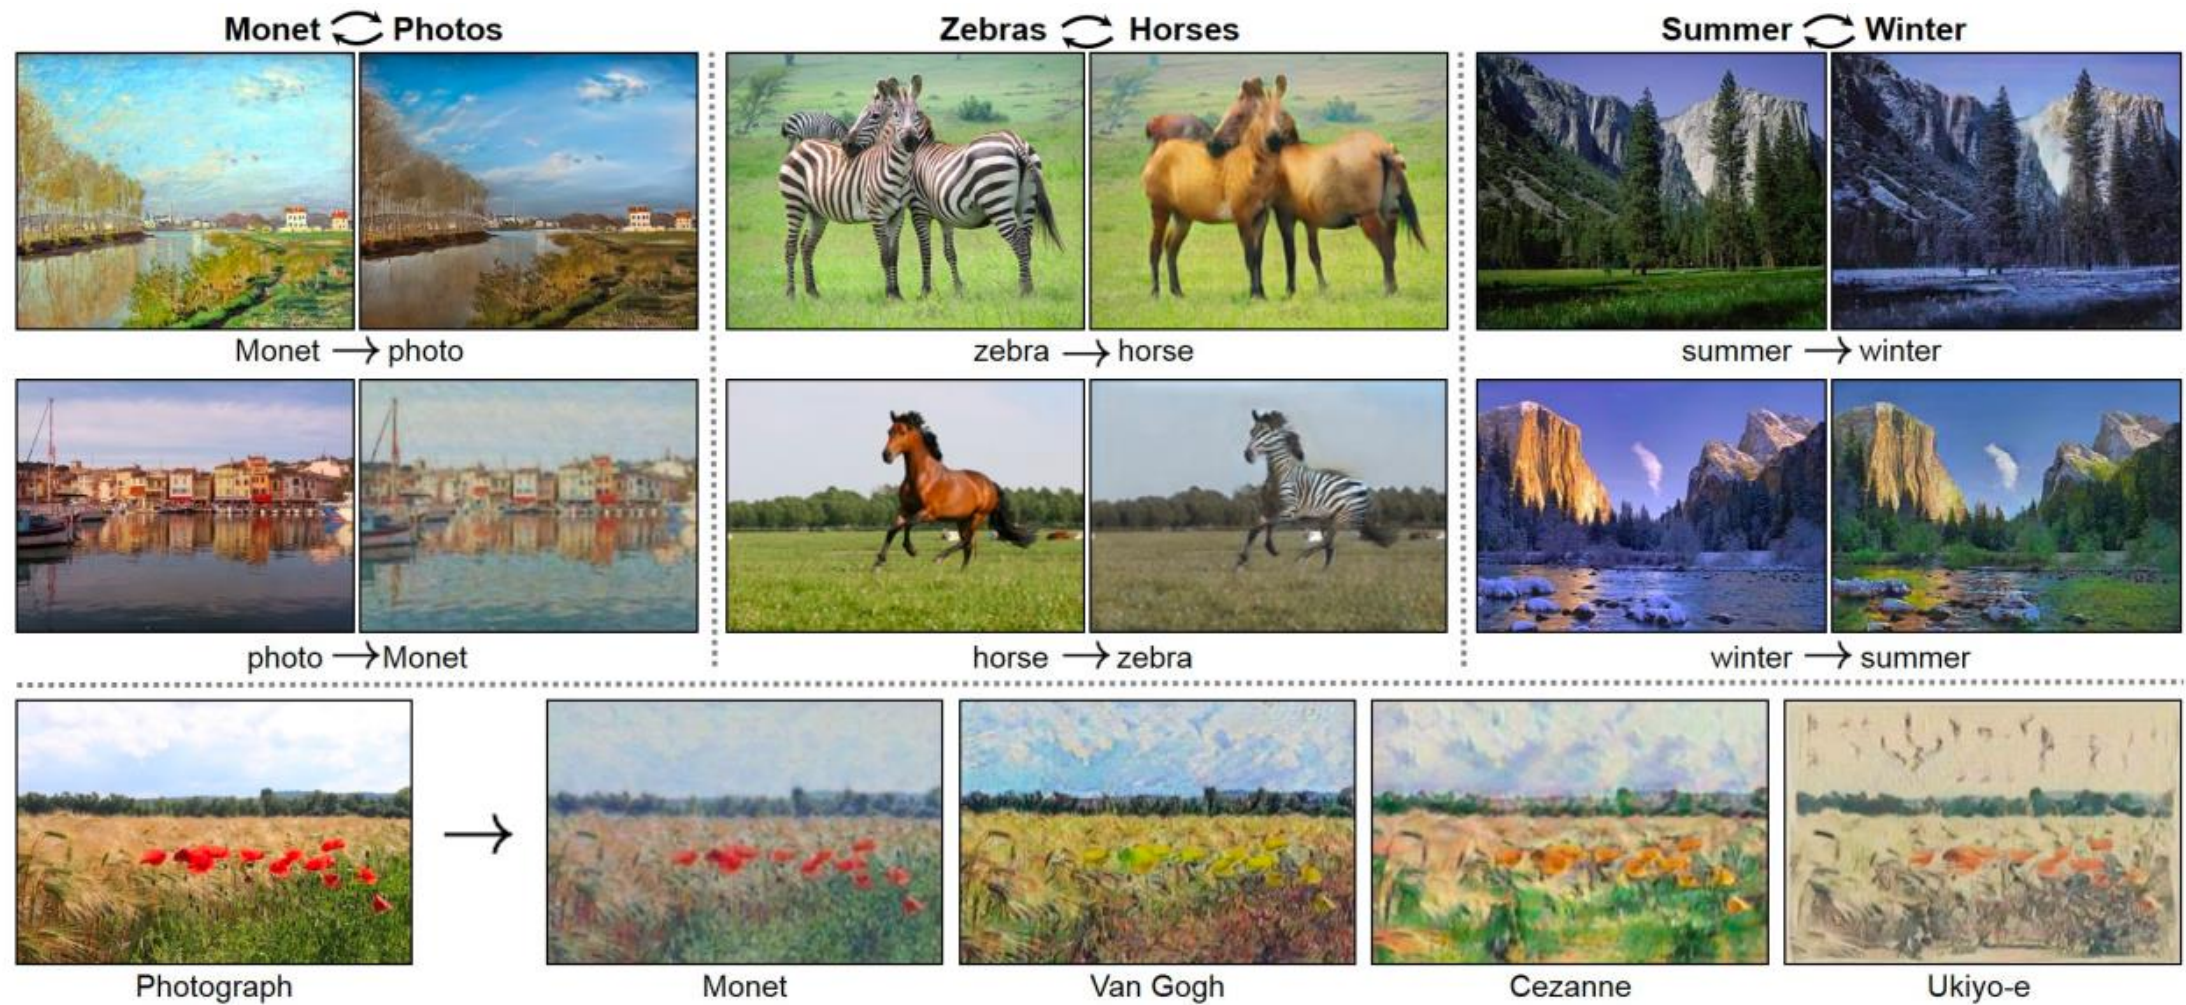
Edges to Photo



Isola et al, "Image-to-Image Translation with Conditional Adversarial Nets", CVPR 2017



# GANs

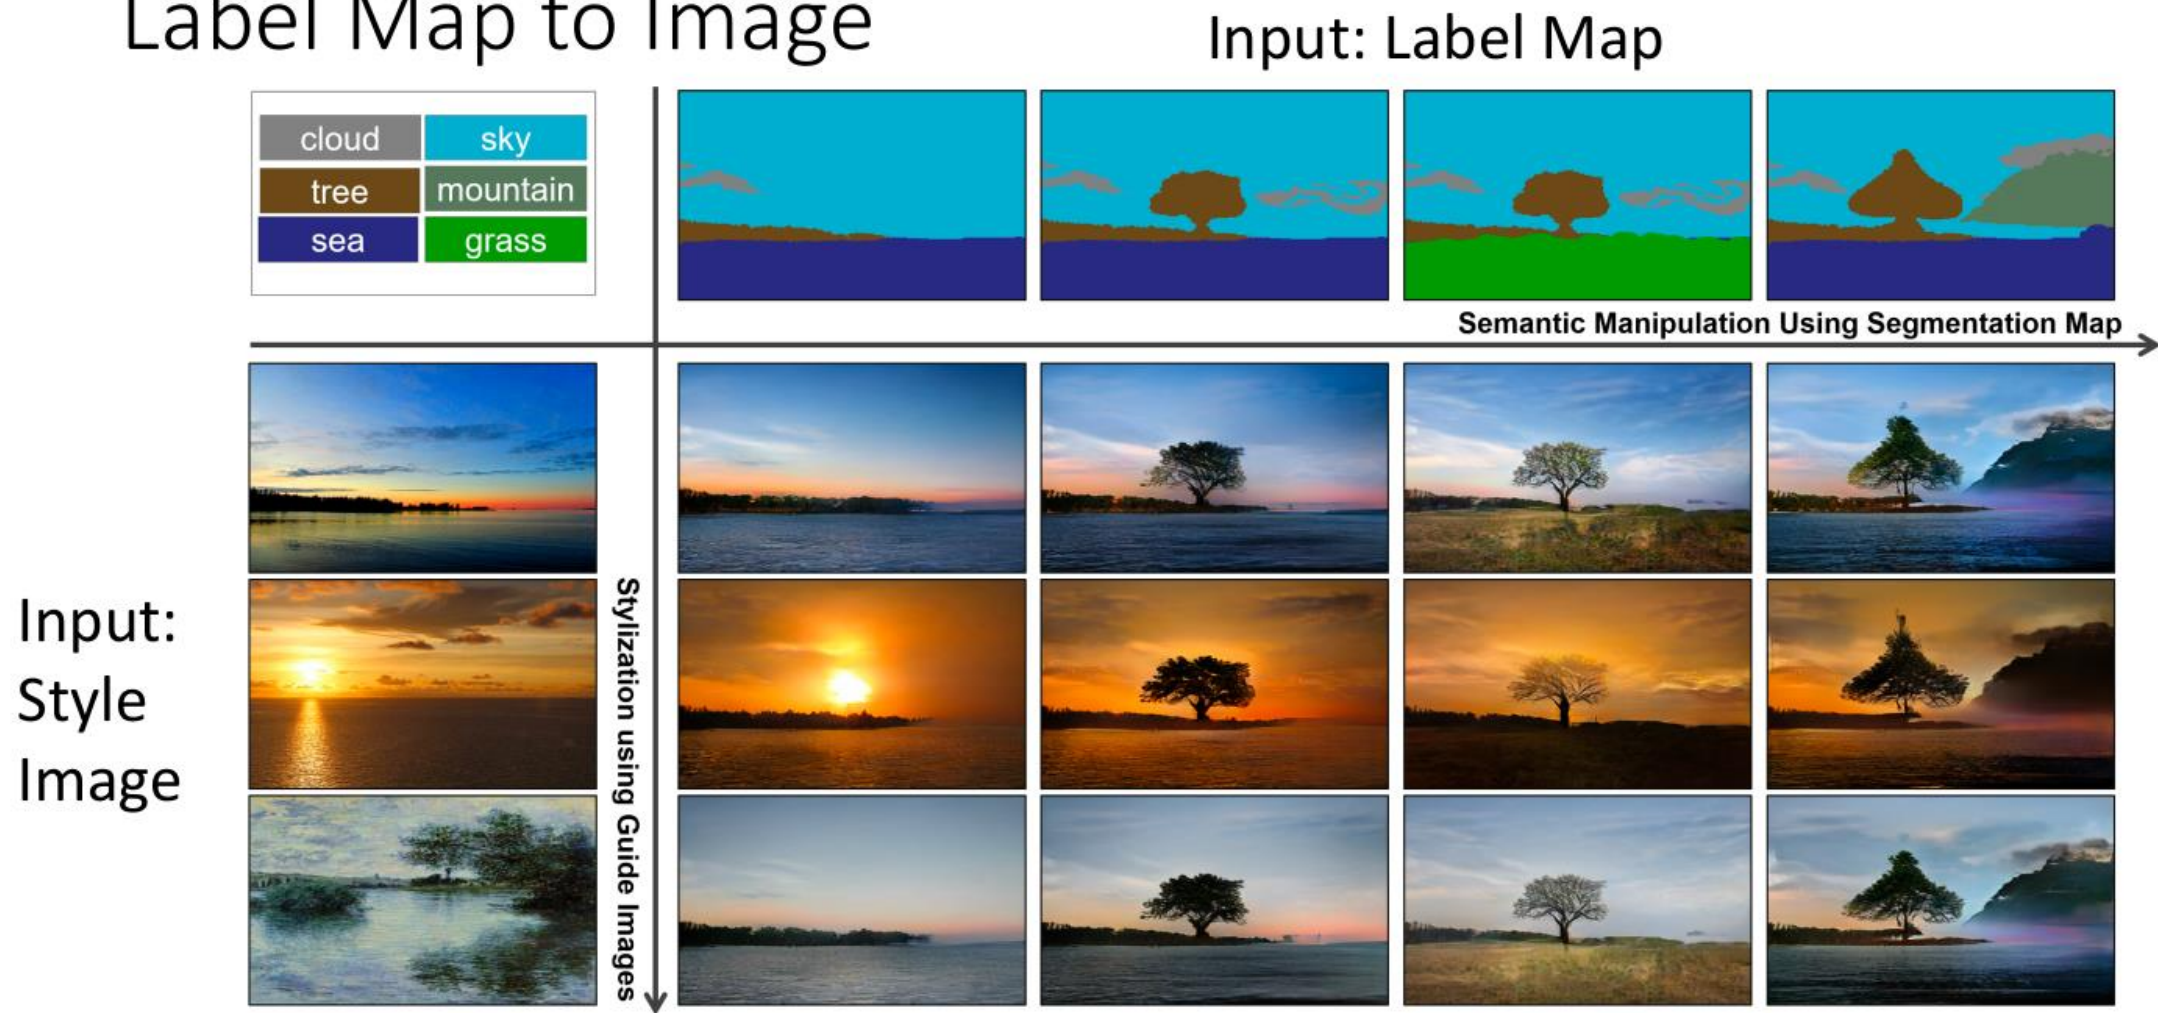


Zhu et al, "Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks", ICCV 2017



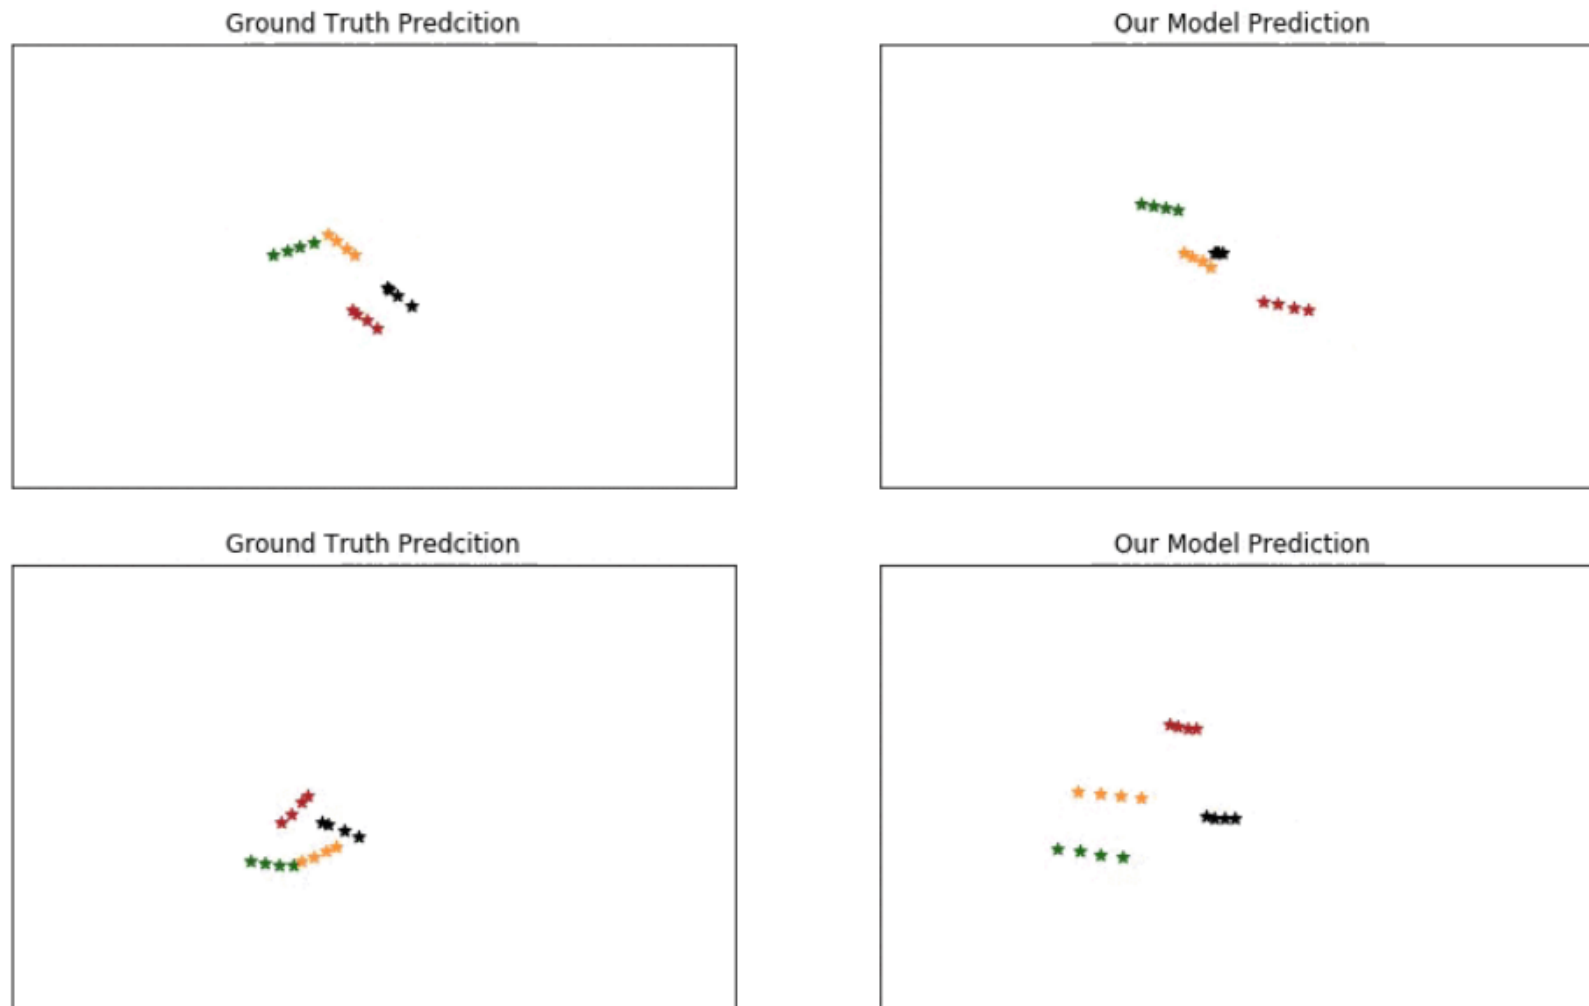
# GANs

## Label Map to Image



Park et al, "Semantic Image Synthesis with Spatially-Adaptive Normalization", CVPR 2019

# GANs



Gupta, **Johnson**, Li, Savarese, Alahi, "Social GAN: Socially Acceptable Trajectories with Generative Adversarial Networks", CVPR 2018

---

**Thank you!**  
**Q & A**